TITLE:

A genotype-phenotype-fitness map reveals local modularity and global pleiotropy of adaptation

AUTHORS:

Grant Kinsler[1*], Kerry Geiler-Samerotte[1,2*], Dmitri Petrov[1]

AFFILIATIONS:

[1]Department of Biology, Stanford University, Stanford, CA

[2]Center of Mechanisms of Evolution, School of Life Sciences, Arizona State University, Tempe, AZ

*Equal contribution

SUMMARY:

Building a genotype to phenotype to fitness map of adaptation is a central goal in evolutionary biology. It is also notoriously difficult even when the adaptive mutations are known because it is hard to identify which phenotypes make these mutations adaptive. We solve this problem by first quantifying how the fitness of hundreds of adaptive mutants responds to subtle environmental shifts and then inferring the existence of fitness-relevant phenotypes implicit in these patterns of fitness variation. We find that a small number of phenotypes predicts the fitness of the adaptive mutations near their original glucose-limited evolution condition. Importantly, phenotypes that matter little to fitness at or near the evolution condition can matter strongly in distant environments. This suggests that adaptive mutants are locally modular — affecting a small number of phenotypes that matter to fitness in the environment where they evolve — yet globally pleiotropic — affecting many phenotypes that contribute to fitness in new environments.

INTRODUCTION

High-replicate laboratory evolution experiments are opening an unprecedented window into the dynamics and genetic basis of adaptive change by *de novo* mutation (Crozat et al., 2010; Good et al., 2017; Lang et al., 2013; Levy et al., 2015; Tenaillon et al., 2012; Venkataram et al., 2016). One of the key insights revealed by these studies is that in many systems, evolution can proceed rapidly via many large-effect single mutations (Venkataram et al., 2016)(Crozat et al., 2010; Good et al., 2017; Lang et al., 2013; Levy et al., 2015; Tenaillon et al., 2012)Venkataram et al., 2016). While the identities of these adaptive mutations are often unique to a specific replicate of the evolutionary experiment, across many replicates they tend to occur in similar genes and pathways. Thus, while the diversity of mutations suggests that there might be many ways to adapt, the much smaller number of apparent functional units implies, in contrast, that most adaptive mutations affect a small set of key phenotypes (Fig 1A).

Consider the seminal study by Tenaillon et al. (Tenaillon et al., 2012) in which 115 populations were evolved at high temperature for ~2000 generations. While the authors identified over a thousand mutations that were largely unique to each population, the number of affected genes was much smaller with 12 genes being hit over 25 times each. Even greater convergence was seen at higher levels of organization such as operons. Similarly, Venkataram et al (Venkataram et al., 2016) find that, of the hundreds of unique genetic mutations that occur during adaptation to glucose-limitation, the vast majority fall into a relatively small number of genes (mostly *IRA1, IRA2, GPB2, PDE2*) and primarily two pathways - Ras/PKA and TOR/Sch9. Thus despite the diversity of mutations, it is possible that all of their effects can be mapped in one or few dimensions required to describe their effects on the Ras/PKA or TOR/Sch9 pathways. These are just two examples, but the pattern has been seen repeatedly (Barghi et al., 2019; Crozat et al., 2010; Good et al., 2017; Lang et al., 2013; Lind et al., 2015). Note that this pattern is seen not only in experimental evolution but also in cancer evolution. Individual tumors are largely unique in terms of specific mutations, but these mutations affect a much smaller set of driver genes and an even smaller number of higher functional units such as signalling pathways (Hanahan and Weinberg, 2011, 2000).

The mapping of adaptive mutations to a smaller number of functional units and thus a low-dimensional space representing their phenotypic effects (Fig 1A) is consistent with theoretical models of adaptation. These theoretical models argue that adaptive mutations, especially those of substantial fitness benefit, cannot affect too many phenotypes at once as most such effects should be deleterious and thus inconsistent with the overall positive effect on fitness (Fisher, 1930; Orr, 2000). More recent studies likewise suggest that selection against mutations with high pleiotropy, *i.e.* mutations that affect many phenotypes, has resulted in a modular architecture of the genotype-phenotype map, in which genetic changes can influence some phenotypes without affecting others (Altenberg, 2005; Collet et al., 2018; Melo et al., 2016; Wagner et al., 2007; Wagner and Altenberg, 1996; Wagner and Zhang, 2011; Welch and Waxman, 2003). This architecture would allow single mutations to drive large low-dimensional phenotypic shifts. It would also explain the observations that very large collections of adaptive mutations are not diverse in terms of affected genes, pathways, and phenotypes. The reason for this is that only mutations that affect the genes, pathways, and phenotypes that represent the correct module most relevant to adaptation in the study environment will be adaptive.

While theoretically appealing, the possibility that observed adaptive mutations indeed affect only a very small number of phenotypes is difficult to reconcile with the notion that organisms are tightly integrated (Kacser and Burns, 1981; Paaby and Rockman, 2013; Rockman, 2012). Further, there is experimental evidence of widespread pleiotropy, for example, from genome wide association studies that suggest every gene can influence every trait, at least to some extent (Boyle et al., 2017; Chesmore et al., 2018; Sella and Barton, 2019; Sivakumaran et al.,

2011; Visscher and Yang, 2016). It is possible that pleiotropy is common, but strongly adaptive mutations observed in experimental evolution are unusual in that they have few pleiotropic phenotypic effects. Another possibility is that these mutations do have pleiotropic side effects but these matter little to fitness in the evolution condition (Fig 1B). Note that here we do not need to claim that these phenotypic effects *never* matter to fitness but rather that they do not matter to fitness in the condition where they evolved. In fact the key prediction of this model is that one should be able to detect such latent pleiotropy by showing that additional phenotypic changes matter to fitness in other environments (Fig 1C).

If the model depicted in Fig 1B and C is true, then it is possible that adaptive mutations are *locally modular* — that they affect very few phenotypes that matter to fitness in the evolution condition — and *globally pleiotropic*. Under this model, the large number of distinct mutations available to adaptation becomes important. Indeed while these mutations tend to influence similar genes and pathways, their phenotypic effects do not simply collapse to a low dimensional functional space. Instead this genetic diversity becomes a source of consequential phenotypic diversity, but only once these genetic variants leave the local environment in which they originated.

In order to test this model and better understand the genotype-phenotype-fitness map, we face the difficult task of identifying which phenotypes adaptive mutations affect and then determining how these phenotypes contribute to fitness. This is a challenging problem as the possible number of phenotypes one can measure is virtually infinite, *e.g.* the expression level of every gene or the quantity of every metabolite (Coombes et al., 2019; Mehlhoff et al., 2020). Further, many measurable phenotypes are related in complex ways (Geiler-Samerotte et al., 2019). Mapping their contribution to fitness requires a complete understanding of how genetic changes lead to molecular changes and how these percolate to higher functional levels and ultimately influence fitness (Kemble et al., 2020). This might be possible to do in some cases where the phenotype to fitness mapping is simple (e.g. antibiotic resistance driven by a specific enzyme or tRNA or protein folding mediating specific RNA or protein function) (Baeza-Centurion et al., 2019; Cowperthwaite et al., 2005; Diss and Lehner, 2018; Domingo et al., 2019; Harmand et al., 2017; Karageorgi et al., 2019; Li and Zhang, 2018; Otwinowski et al., 2018; Pressman et al., 2019; Sarkisyan et al., 2016; Starr et al., 2018; Weinreich, 2006) but is exceptionally difficult for complex phenotypes.

Moreover, to distinguish the model in Fig. 1A from Fig 1B, we need to understand these genotype-phenotype-fitness maps not only in the environment in which adaptive mutants evolved, but also in other environments, as depicted in Fig. 1C. And we need to do this for many adaptive mutants so that we can assess the extent to which different mutants affect different phenotypes. Considering the scope of this challenge, it is not surprising that despite much theoretical discussion of modularity and pleiotropy as it relates to adaptation, experimental approaches to address these questions have lagged behind.

Here we suggest a way to model the genotype-phenotype-fitness relationship that avoids the problem of measuring each phenotype and its effect on fitness explicitly. We argue that it is possible to investigate the genotype-phenotype-fitness map by comparing how the fitness effects of many mutations change across a large number of environments. The way each mutant's fitness varies across environments must be related to its phenotype, and thus the way mutants co-vary in fitness across environments tells us whether they affect similar fitness-relevant phenotypes. We can use the profiles of fitness across a set of environments to identify the total number of fitness-relevant phenotypes affected across a collection of adaptive mutants, the extent to which different mutants affect different phenotypes, and whether the contribution of each phenotype to fitness changes across environments.

Here we build a genotype-phenotype-fitness model for hundreds of adaptive mutations that originally evolved in a glucose-limited environment. We use this model to accurately predict the fitness of these mutants across a set of 45 environments that vary in their similarity to the evolution condition. We find that the behavior of adaptive mutations can be described by a low-dimensional phenotypic model. In other words, these mutants affect a small number of phenotypes that matter to fitness in the glucose-limited condition in which they evolved. We find that this low-dimensional phenotypic model makes accurate predictions of mutant fitness in novel environments even when they are distant from the evolution condition. Moreover, we find that some phenotypes that contribute very little to fitness in the evolution condition become surprisingly important in some novel environments. This suggests that adaptive mutations are globally pleiotropic in that they affect many phenotypes overall, but that they are locally modular in that only a small number of these phenotypes have substantial effects on fitness in the environment they evolved in. Overall, we suggest that this set of adaptive mutations contains substantial and consequential latent phenotypic diversity, meaning that despite targeting similar genes and pathways, different adaptive mutants may respond differently to future evolutionary challenges. This finding has important consequences for understanding how directional selection can generate consequential phenotypic heterogeneity both in natural populations and also in the context of diseases such as cancer and viral or bacterial infections. In addition, our results show that our abstract, top-down approach is a promising route of analysis for investigating the phenotypic and fitness consequences of mutation.

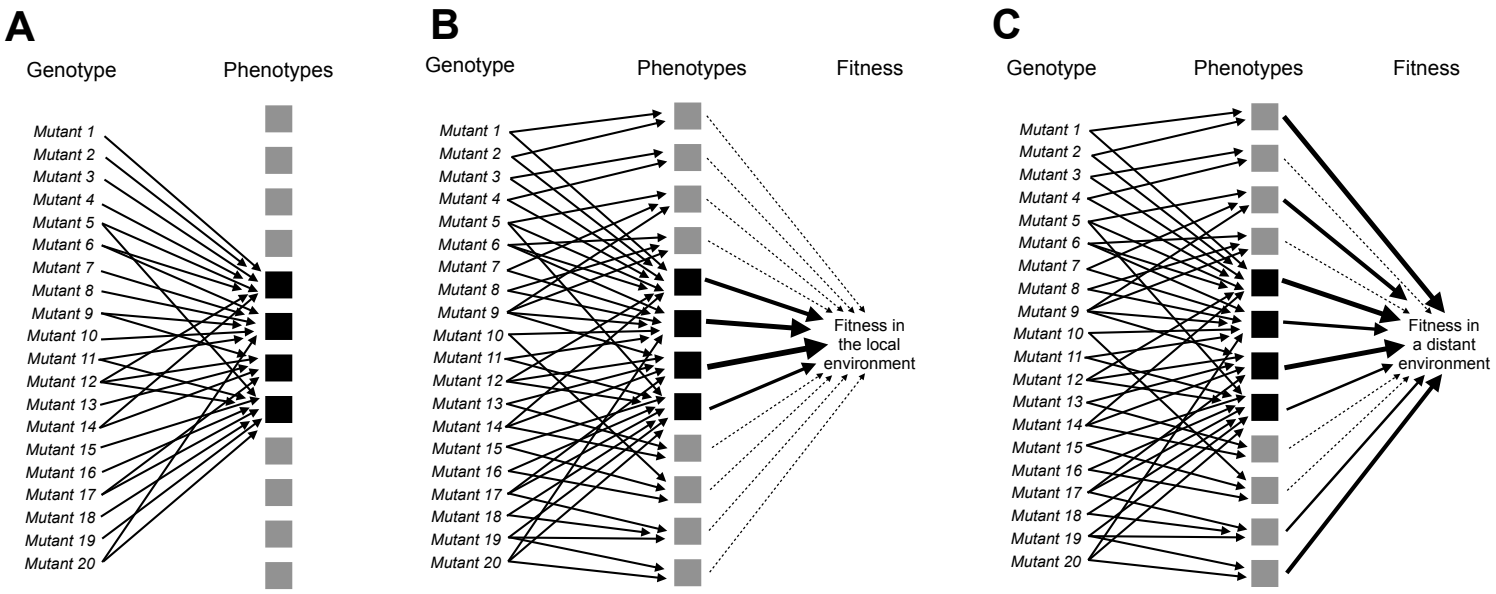

Figure 1. Adaptive mutations can be locally modular and globally pleiotropic. (A) A collection of adaptive mutations may affect a small number of phenotypes (four black squares). **(B)** Alternatively, these mutations may collectively (and individually) affect many phenotypes, but only a small number of phenotypes may matter to fitness (those indicated by black squares with thick arrows pointing to fitness), whereas the other phenotypes may make very small contributions to fitness (those indicated by the gray squares and thin, dashed lines leading to fitness). **(C)** Under the model in B, the contribution of each phenotype to fitness can change depending on the environment. Thus fitness differences between seemingly similar mutants can be revealed by measuring fitness in more environments. Such fitness differences suggest the presence of phenotypic differences between mutants.

RESULTS

Mutants that improve fitness under glucose limitation vary in their genotype-by-environment interactions

A previous evolution experiment generated a collection of hundreds of independent mutations that each provide a benefit to yeast cells growing in a glucose-limited environment (Levy et al., 2015). Many of these mutants, which began the evolution experiment as haploids, underwent whole-genome duplication to become diploid, which improved their relative fitness (Venkataram et al 2016). Some of these diploids acquired additional mutations, including amplifications of either chromosome 11 or 12 as well as point mutations, which generated additional fitness benefits. The adaptive mutants that remained haploid acquired both gain- and loss-of-function mutations in nutrient-response pathways (Ras/PKA and TOR/Sch9). Some other mutations were also observed, including a mutation in the HOG pathway gene *SSK2* (Venkataram et al., 2016). Although these mutants have been well-characterized at the level of genotype and fitness, it is unclear what phenotypes they affect. The first question we address is whether these diverse mutations collectively affect a large number of phenotypes that matter to fitness, or whether these mutants are functionally similar in that they collectively alter a small set of fitness-relevant phenotypes.

Understanding the map from genotype to phenotype to fitness is extremely challenging because each genetic change can influence multiple traits, not all of which are independent or contribute to fitness in a meaningful way. We contend with this challenge by measuring how the relative fitness of each adaptive mutant changes across a large collection of similar and dissimilar environments, which we term the “fitness profile”. When a group of mutants demonstrate similar responses to environmental change, we conclude that these mutants affect similar phenotypes. By clustering mutants with similar fitness profiles across a collection of environments, we can learn about which mutants influence similar phenotypes, as well as estimate the total number of fitness-relevant phenotypes represented across all mutants and all investigated environments.

Because our mutant strains are barcoded, we can use previously-established methods to measure their relative fitness in bulk and with high-precision (Venkataram et al., 2016). Specifically, we compete a pool of the barcoded mutants against an ancestral reference strain over the course of several serial dilution cycles. During each 48 hour cycle, the yeast are given fresh glucose-limited media which supports 8 generations of exponential growth after which glucose is depleted and cells transition to non-fermentable carbon sources. After every 48 hour cycle, we transfer ~5x10^7 cells to fresh media to continue the growth competition. We also extract DNA from the remaining cells to PCR amplify and sequence their barcodes. We repeat this process four times, giving us an estimate of the frequency of each barcode at five time-points. By quantifying the log-linear changes in each barcode’s frequency over time and correcting for the mean-fitness change of population, we can calculate the fitness of each barcoded mutant relative to the reference strain (Fig 2A; Methods).

Using this method, we quantify the fitness of a large number of adaptive mutants in 45 environments. We focus on a set of 292 adaptive mutants that have been sequenced, show clear adaptive effects in the glucose-limited condition in which these mutants evolved (hereafter “evolution condition”; EC) (Fig 2B; Table S1), and for which we obtained high-precision fitness measurements in all 45 environments. These environments include some experiments from previously published work (Li et al., 2018; Venkataram et al., 2016), as well as 32 new environments including replicates of the evolution condition, subtle shifts to the amount of glucose, changes to the shape of the culturing flask, changes to the carbon source, and addition of stressors such as drugs or high salt (Table S2).

In order to determine the total number of phenotypes that are relevant to fitness in the EC, we focus on environments that are very similar to the EC but still induce small yet detectable perturbations in fitness. We do so because the phenotypes that are the most relevant to fitness may change with the environment (Fig 1B and Fig 1C). Thus, we partition the 45 environments into a set of “subtle” perturbations, from which we will detect the phenotypes relevant to fitness

near the EC, and “strong” perturbations which we will use to study whether these mutants influence additional phenotypes that matter in other environments (Fig 1C).

To partition environments into subtle and strong perturbations of the EC, we rely on the nested structure of replicate experiments performed in the EC. We performed nine such replicates, each at different times, which each included multiple replicates performed at the same time. We observe much less variation across replicates performed simultaneously than across replicates performed at different times ($p < 1e\text{-}5$ from permutation test). Variation across experiments performed at different times is often referred to as “batch effects” and likely reflects environmental variability that we were unable to control (e.g. slight fluctuations in incubation temperature due to limits on the precision of the instrument). These environmental differences between batches are as subtle as possible, as they represent the limit of our ability to minimize environmental variation. Thus, variation in fitness across the EC batches serves as a natural benchmark for the strength of environmental perturbations. If the deviations in fitness caused by an environmental perturbation are substantially stronger than those observed across the EC batches, we call that perturbation “strong”.

More explicitly, to determine whether a given environmental perturbation is subtle or strong, we subtract the fitness of adaptive mutants in this environment from their average across the EC batches. We then compare this difference to the variation in fitness observed across the EC batches. Sixteen environmental perturbations provoked fitness differences that were similar to those observed across EC batches (Z-score < 2). These environments, together with the nine EC batches, make up a set of subtle environmental perturbations. The remaining 20 environments, where the average deviation in fitness is substantially larger than that observed across batches (Z-score > 2), were classified as strong environmental perturbations (Fig 2C, top; Methods).

The rank order of the fitnesses of many mutations is largely preserved across the 25 environments that represent subtle perturbations (Fig 2C, bottom). For example, *IRA1* nonsense mutants, which are the most adaptive in the EC, generally remain the most adaptive across the subtle perturbations. Additionally, the *GPB2* and *PDE2* mutants have similar fitness effects across EC batches and only occasionally switch order across the subtle environmental perturbations. In contrast, the 20 environments that represent strong perturbations reveal clear genotype-by-environment interactions (Fig 2C, bottom). For example, altering the transfer time from 48 to 24 hours (the “1 Day” environment in Fig 2C) affects *GPB2* mutants more strongly compared to the other mutants in the Ras/PKA pathway, including *IRA1* and *PDE2*. The strongest environmental perturbations reveal clear tradeoffs for some of these adaptive mutants. For example, *PDE2* and *IRA1-nonsense* but not *GPB2* mutants are particularly sensitive to osmotic stress as indicated by the NaCl and KCl environments. Additionally, *IRA1-nonsense* mutants become strongly deleterious in the long transfer conditions that experience stationary phase (5-, 6-, 7-Day environments) (Li et al., 2018). In contrast to complex behavior exhibited by the adaptive haploids, the diploids appear to be relatively robust to strong tradeoffs, appearing similarly adaptive across all perturbations, subtle and strong.

The observation that different mutants have different and fairly complex fitness profiles suggests that they have different phenotypic effects. Even *PDE2* and *GPB2,* which have similar fitnesses in the EC and are negative regulators of the same signalling pathway, have different fitness profiles. Do these diverse phenotypic effects contribute to fitness in the EC? To examine how many phenotypes matter to fitness in the EC, we test whether it is possible to create low dimensional models that capture the complexity of the fitness profiles of all adaptive mutants across all subtle perturbations.

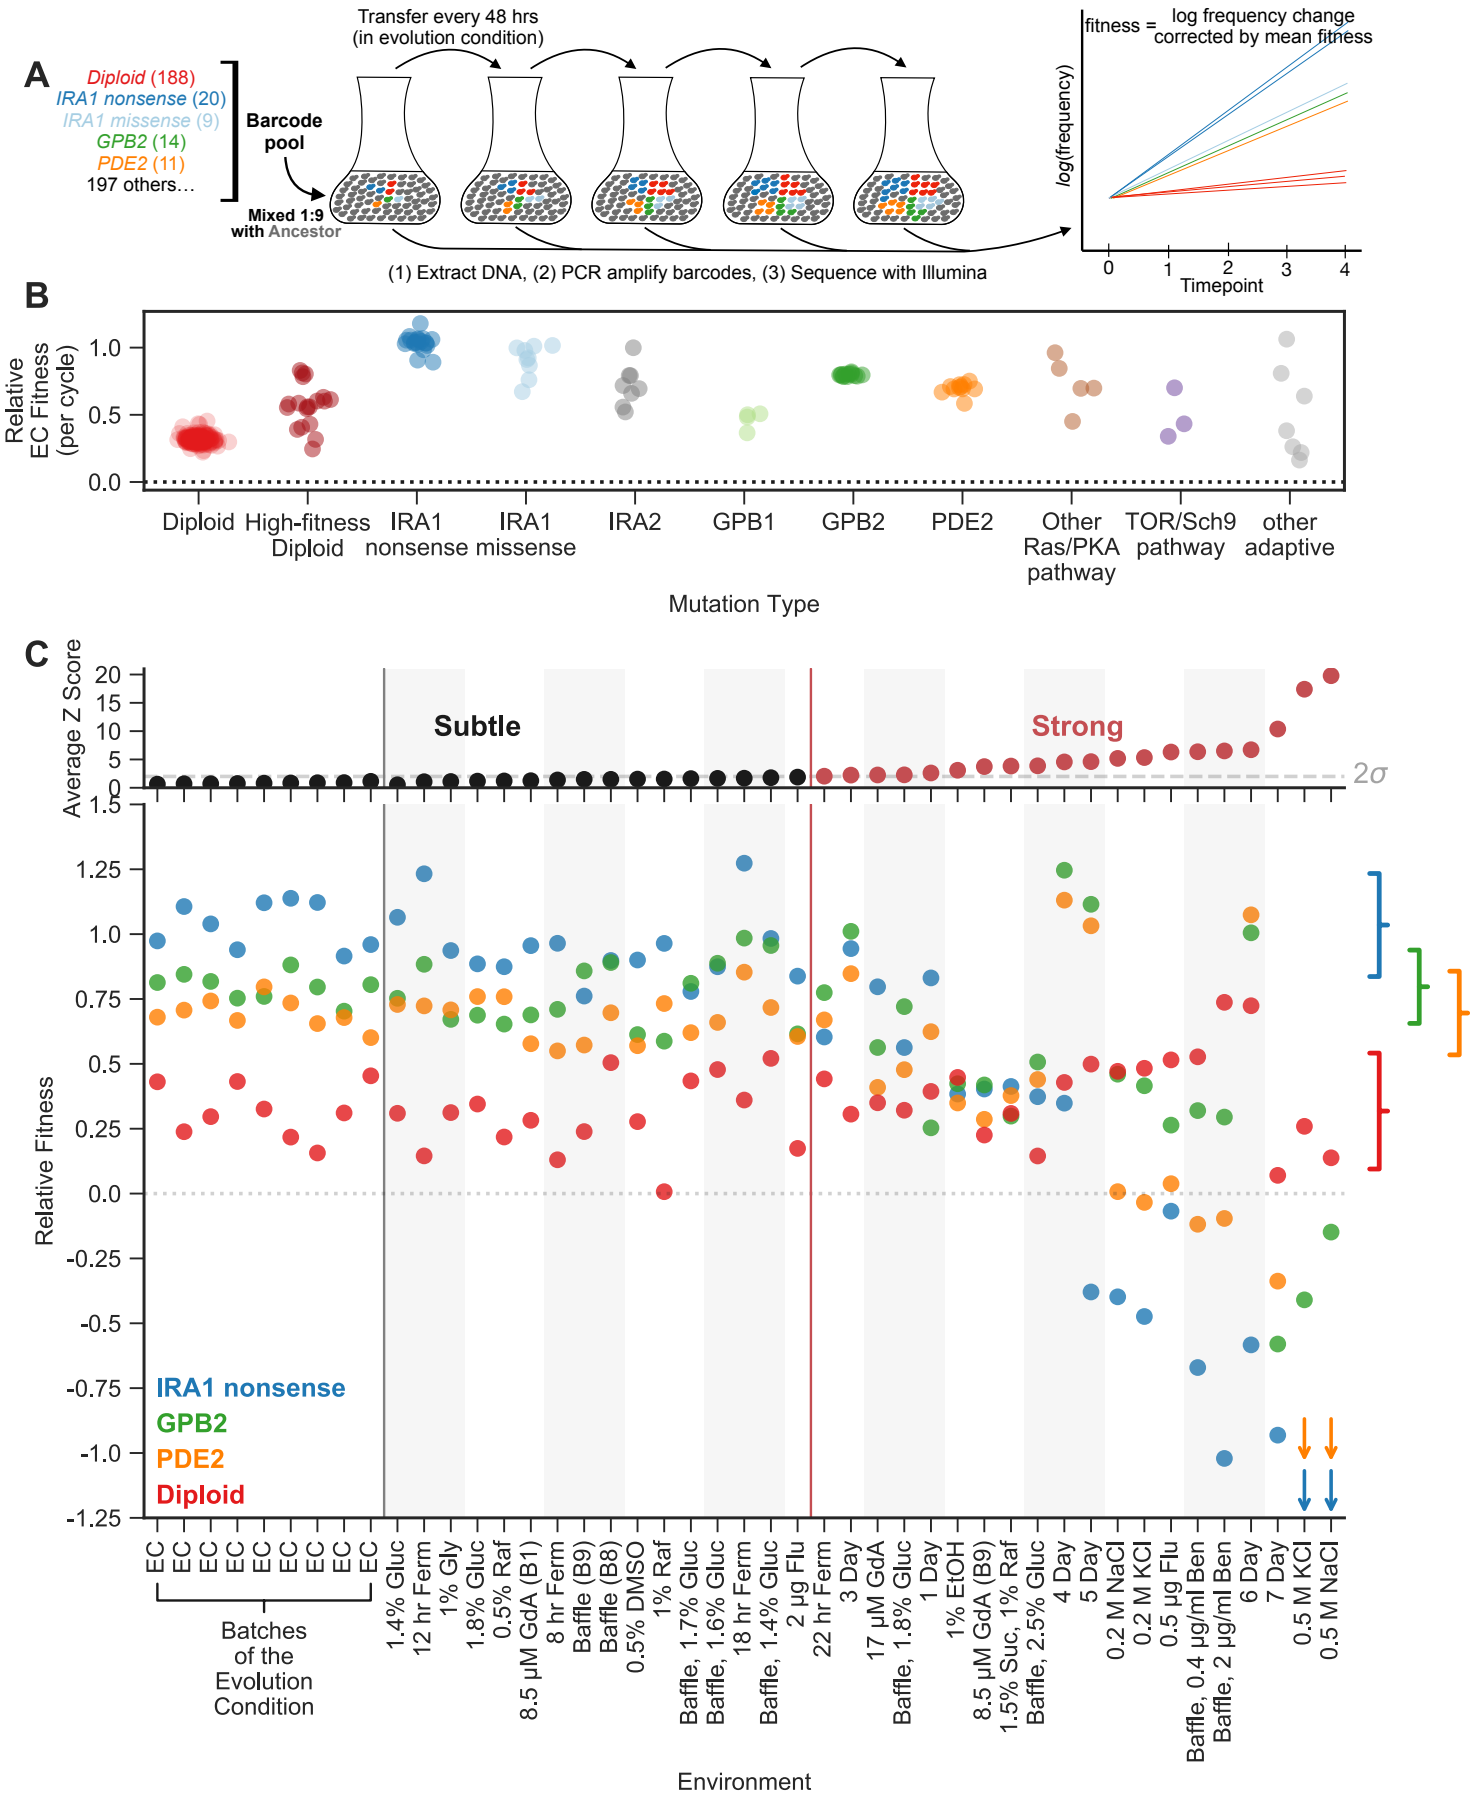
A
Diploid (188)
IRA1 nonsense (20)
IRA1 missense (9)
GPB2 (14)
PDE2 (11)
197 others…
Barcode pool
Mixed 1:9 with Ancestor
Transfer every 48 hrs (in evolution condition)
(1) Extract DNA, (2) PCR amplify barcodes, (3) Sequence with Illumina
fitness = log frequency change corrected by mean fitness
log(frequency)
Timepoint
B
Relative EC Fitness (per cycle)
Diploid
High-fitness Diploid
IRA1 nonsense
IRA1 missense
IRA2
GPB1
GPB2
PDE2
Other Ras/PKA pathway
TOR/Sch9 pathway
other adaptive
Mutation Type
C
Average Z Score
Subtle
Strong
2σ
Relative Fitness
IRA1 nonsense
GPB2
PDE2
Diploid
Batches of the Evolution Condition
EC
1.4% Gluc
12 hr Ferm
1% Gly
1.8% Gluc
0.5% Raf
8.5 µM GdA (B1)
8 hr Ferm
Baffle (B9)
Baffle (B8)
0.5% DMSO
1% Raf
Baffle, 1.7% Gluc
Baffle, 1.6% Gluc
18 hr Ferm
Baffle, 1.4% Gluc
2 µg Flu
22 hr Ferm
3 Day
17 µM GdA
Baffle, 1.8% Gluc
1 Day
1% EtOH
8.5 µM GdA (B9)
1.5% Suc, 1% Raf
Baffle, 2.5% Gluc
4 Day
5 Day
0.2 M NaCl
0.2 M KCl
0.5 µg Flu
Baffle, 0.4 µg/ml Ben
Baffle, 2 µg/ml Ben
6 Day
7 Day
0.5 M KCl
0.5 M NaCl
Environment

Figure 2. Measuring fitness for a collection of adaptive mutants across many environments reveals gene-by-environment interactions. (A) Schematic of fitness measurement procedure. Adaptive mutants tagged with DNA barcodes are pooled at a 1:9 ratio with an ancestral reference strain. The pool is then propagated for several growth cycles, where the population is diluted into fresh media at fixed time intervals. DNA is extracted from each time-point, and the barcode region is PCR amplified and then sequenced. A mutant's relative fitness is calculated based on the rate of change of its barcode's frequency, corrected for the mean fitness of the population (see Methods). Relative fitness is calculated in units of "per cycle", representing the improvement of each barcode relative to the reference over the course of the time between transfers. **(B)** Relative fitness of each mutant in the evolution condition, calculated as the average across all 9 Evolution Condition (EC) batches. **(C) (top)** Environments are ordered from left to right depending on the degree to which they perturb mutant fitness from the average fitness observed across all EC batches. Environments in which average mutant fitness is within two standard deviations of average mutant fitness across EC batches are denoted in black and make up the subtle perturbation set. Environments in which aggregate mutant behavior exceeds two standard deviations are shown in red and make up the strong perturbations set. **(bottom)** This plot displays, for the four most common types of adaptive mutation observed in response to glucose limitation (Venkataram et al., 2016), the average fitness in each of the 45 environments we study. Brackets on the right represent the amount of variation in fitness observed for each type of mutation across the EC batches, with the notch representing the mean and the arms representing two standard deviations on either side of the mean.

A model including 8 fitness-relevant phenotypes captures fitness variation across subtle environmental perturbations

We utilize these complex fitness profiles to estimate the number of phenotypes that contribute to fitness in the EC. Given that many of these mutants affect genes in the same nutrient response pathway, the number of unique phenotypes they affect may be small. Alternatively, given the observation that these mutants have different interactions with environments that represent strong perturbations (Fig 2C), this number may be large. We use singular value decomposition (SVD) to ask how much of the complexity in these fitness profiles can be captured by a low dimensional phenotypic model (Fig 3A). SVD creates a model, consisting of two abstract multi-dimensional spaces.

The first space, *P*, represents the phenotypic effects of each mutant as a dimension (there are *k* phenotypic dimensions depicted in Fig 3A). Each mutant is represented by coordinates specifying a location in the phenotype space *P* (e.g. mutant 1 having coordinates$(p_{11}, p_{12}, p_{13}, \ldots, p_{1k})$). The ancestral reference lineage, which, by definition, has relative fitness zero in every environment, is placed at the origin (e.g. (0, 0, 0, … 0)) in this phenotypic space. In this sense, we can think of a mutation's effect on any phenotype as a measure of the distance from the location of the mutant to the origin.

The second space, *E*, represents the contribution of each of the phenotypes in *P* to fitness, and thus has the same number of dimensions as *P*. If a phenotype does not contribute substantially to fitness in any environment, it is not represented as a dimension in either space. Therefore, our model captures only fitness-relevant phenotypes. In space *E*, each environment is represented by coordinates specifying a location (e.g. environment 1 having coordinates $(e_{11}, e_{12}, e_{13}, \ldots, e_{1k})$). These coordinates in *E* reflect the contribution (weight) of each of the *k* phenotypic dimensions on fitness in that environment. For example, an environment where only a single phenotype matters to fitness would be placed at the origin for all the axes, except for the axis corresponding to the single phenotypic dimension that matters. Environments for which the same phenotypes contribute to fitness will be placed closer together in the space *E*.

In this model, each phenotype contributes to fitness independently, by definition, such that the fitness of mutant *i* in environment *j* is determined by each phenotypic effect of mutant *i,*scaled by

the contribution of that phenotype to fitness in environment *j*. A linear combination of these weighted phenotypic effects determines the fitness of mutant *i* in environment *j*:

$$f_{ij} = p_{i1}e_{1j} + p_{i2}e_{2j} + p_{i3}e_{3j} + \ldots + p_{ik}e_{kj}$$

In this model, mutants with similar fitness profiles, for example mutants 1 and 2 in Fig 3A, will be inferred as having similar phenotypic effects, and thus be located near each other in the phenotypic space *P*. Mutants with dissimilar fitness profiles, for example mutants 3 and 4 in Fig 3A, can be inferred to have at least some differing phenotypic effects, which might be mediated by a different effect on a single phenotypic component or different effects on many. Mutants with dissimilar fitness profiles are informative about the number of dimensions needed in this abstract model of phenotypic space.

Here, we focus on counting the number of phenotypes that contribute to fitness in the original glucose-limited environment in which these adaptive mutants evolved. We used SVD to build an abstract model that captures fitness profiles of all 292 adaptive mutants across the 25 subtle perturbations. This model suggests that the majority of the variation in fitness for the 292 adaptive mutants across the 25 subtle perturbations can be explained by eight phenotypic dimensions. The first phenotypic component is very large and explains 95% of variation in fitness across all mutants and all subtle perturbations (Fig 3A). This component captures the variation in fitness explainable in the absence of genotype-by-environment interactions, where each mutation has a single effect that is scaled by the environment. As such, this first component effectively represents each mutant's average fitness in the EC (Fig S2) and the average impact of each subtle perturbation on mutant fitness (Fig S2). It is not surprising that this component explains much of this variation, as the fitness of mutants in the EC should be predictive of fitness in similar environments. The next seven components capture additional variation not detectable from the simple 1-component model and thus represent genotype-by-environment interactions. Of these, the first four capture 87% of the variation not captured by component one (67.8%, 8.3%, 5.6%, and 5.3%, respectively). The remaining three interaction components each capture less than 2% of the variation not captured by component one (Fig 3A). We cannot distinguish any additional components, beyond these eight, from noise. This is because we see components that explain a similar amount of variation when we apply SVD to datasets composed exclusively of values generated by our noise model (Fig3A; see Methods for additional details).

We confirm that these eight phenotypic components capture meaningful biological variation in fitness using bi-cross-validation. Specifically, we designate a balanced set of 60 of the 292 mutants as a training set, chosen such that the recurrent mutation types — diploids, high-fitness diploids, Ras/PKA mutants — are roughly equally represented (see Methods). The remaining 232 mutants comprise the test set. This set contains all mutation types represented by only a single mutant, including all TOR/Sch9 (*TOR1, SCH9, KOG1)* and HOG (*SSK2)* pathway representatives, as well as the rest of the recurrent mutants that were not picked for the training set. We iteratively construct phenotype spaces using the 60 training mutants while holding out one subtle perturbation at a time and creating the space with the data from the remaining 24 subtle perturbations. We then predict the fitness of the 232 held-out testing mutants in the held-out condition. We do so using all 8 components, and again with only 7, 6, and so on. Then, we ask whether the 8 component model does a better job at predicting mutant fitness than the other, lower dimensional models. If a component reflects measurement noise rather than biological signal, then the inclusion of this component would lead to overfitting and should harm the model's ability to predict fitness in the held-out data. Instead we find that, on average across the 25 iterations, prediction power improves from the inclusion of each of the eight components. This confirms that even the smallest of these components captures biologically meaningful variation in fitness across the 25 subtle perturbations of the EC. However, the gain in predictive power decreases for each component. The model with only the first component explains on

average 85% of weighted variance for the test mutants in the left-out conditions. A model with only the top five components explains 95.1%, and all eight components explain 96.2% of variation. This suggests that the last few components have very small contributions to fitness in the environments near the EC.

A model including 8 fitness-relevant phenotypes recapitulates known features of adaptive mutations

We next ask whether the 8-dimensional phenotypic space clusters adaptive mutants found in similar genes or pathways (e.g. Ras/PKA or TOR/Sch9), or that represent similar mutation types (haploid v. diploid). We use Uniform Manifold Approximation and Projection (UMAP) to visualize the distance between all the mutants in this phenotypic space. Since the first phenotypic dimension captures the average fitness of each mutant in the EC, and since we already know that mutations to the same gene have similar fitness in the EC (Fig. 2B), we exclude the first phenotypic dimension from this analysis, though the inclusion of the first component does not change the identity of the clusters (Fig S3). By focusing on the other 7 components, we are asking whether genotype-by-environment interactions also cluster the mutants by gene, mutation type, and pathway.

These 7 genotype-by-environment interactions indeed cluster the adaptive mutants by type and by gene (Fig 3B). Specifically, the diploids, *IRA1-nonsense, GPB2,* and *PDE2* mutants each form distinct clusters ($p = 0.0001$, $p = 0.006$, $p = 0.0001$, and $p = 0.0001$, respectively). Interestingly, six high-fitness diploids (diploids with higher than average diploid fitness in the EC) also form a distinct cluster ($p = 0.0001$) despite whole genome sequencing having revealed no mutations in their coding sequences (Fig. 3B). To generate p-values, we calculated the median pairwise distance, finding that multiple mutations in the same cluster are indeed more closely clustered than randomly chosen groups of mutants.

Interestingly, the three smallest components, which capture very little variation in fitness across the environments that reflect subtle perturbations of the EC, cluster some mutants by genotype (Fig S3). Specifically, *PDE2*, *GPB2*, and *IRA1-nonsense* mutants are each closer to mutants of their own type than to other adaptive haploids ($p = 0.0001$, $p = 0.0001$, and $p = 0.03$, respectively). Note that the space defined by the three smallest components does not cluster *IRA1-nonsense* mutants away from diploids ($p = 0.718$). This suggests that some mutants, e.g. *IRA1-nonsense* and diploids, have smaller effects on these three phenotypic components. Overall, our abstract phenotypic model, which reflects the way that each mutant's fitness changes across environments, reveals that mutations to the same gene tend to interact similarly with the environment. This indicates that our approach is a useful and unbiased way to identify mutations that share functional effects (Li et al., 2018).

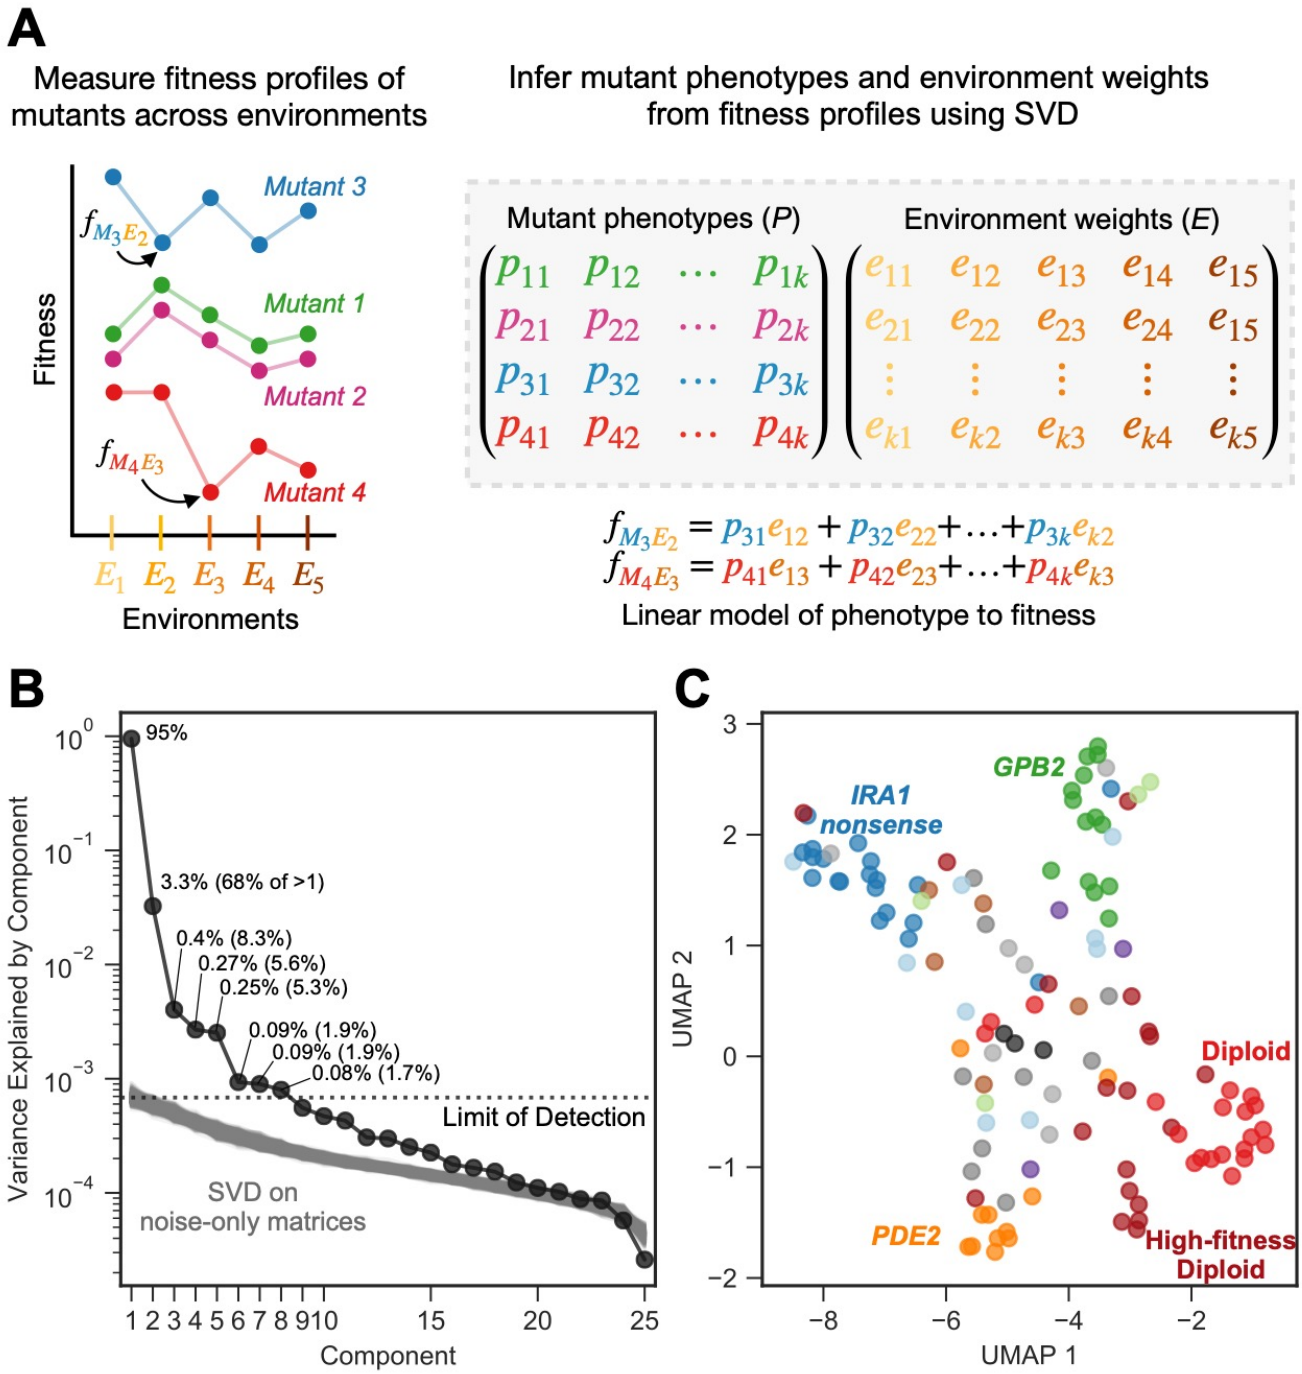

Figure 3. Subtle environmental perturbations reveal an 8-component phenotypic model that reflects known biological features. (A) To infer fitness-relevant phenotypes, we measure the fitness of mutants in a collection of environments and compare their fitness profiles. Mutants with similar fitness profiles (mutants 1 and 2) are inferred to have similar effects on phenotypes. Mutants with dissimilar fitness profiles (mutants 3 and 4) are inferred to have dissimilar phenotypic effects. We use SVD to decompose these fitness profiles into a model consisting of two abstract spaces: one that represents the fitness-relevant phenotypes affected by mutants (*P)* and another which represents the degree to which each phenotype impacts fitness in each environment (*E*). Here, we represent the model with *k* fitness-relevant phenotypes. The model's estimate for fitness for a particular mutant in a particular environment is a linear combination of each mutant phenotype (mutant 1 is represented by the vector $(p_{11}, p_{12}, p_{13}, \ldots, p_{1k})$) scaled by the degree to which that phenotype affects fitness in the relevant environment (environment 1 is represented by the vector$(e_{11}, e_{12}, e_{13}, \ldots, e_{1k})$). We show two examples of the equation used to estimate fitness for the mutants and environments highlighted in the left panel. Note that, for presentation purposes, we show SVD as inferring two matrices. It in fact infers three, but is consistent with our presentation if you fold the third matrix, which represents the singular values, into *E* (see Methods). **(B)** Decomposing the fitness profiles of 292 adaptive mutants across 25 subtle

environmental perturbations reveals 8 fitness-relevant phenotypic components. The variance explained by each component is indicated as a percentage of the total variance. The percentages in parentheses indicate the relative amount of variation explained by each component when excluding the first component. Each of these components explain more variation in fitness than do components that capture variation across a simulated dataset in which fitness varies due to measurement noise. These simulations were repeated 1000 times (grey lines) and used to define the limit of detection (dotted line). **(C)** An abstract space containing 8 fitness-relevant phenotypic components reflects known biological features. This plot shows the relationships of the mutants in a 7-dimensional phenotypic space that excludes the first component, visualized using Uniform Manifold Approximation and Projection (UMAP). Mutants that are close together have similar fitness profiles and are inferred to have similar effects on fitness-relevant phenotypes. Mutants with mutations in the same gene tend to be closer together than random, in particular *IRA1 nonsense* mutants in dark blue, *GPB2* mutants in dark green, *PDE2* mutants in dark orange, and diploid mutants in red. Six high-fitness diploid mutants also form a cluster despite no known genetic similarities.

Fitness variation across subtly different environments predicts fitness in substantially different environments

Now that we have identified the phenotypic components that contribute to fitness in environments that represent subtle perturbations of the EC, we can test the ability of these phenotypic components to predict fitness in more distant environments. Specifically, we can measure how the contribution of each of these components to fitness changes in new environments. We can also determine whether the phenotypic components that contribute very little to explaining fitness variation near the EC might at times have large explanatory power in distant environments (as depicted in Fig 1B and 1C).

To test this we performed bi-cross-validation, using the eight component model constructed from fitness variation of 60 training mutants across 25 subtly different environments to predict the fitness of 232 test mutants in the environments that represent strong perturbations of the EC. To evaluate the predictive power of the model, we compare our model's fitness predictions in each environment to predictions made using the average fitness in that environment. Thus, negative prediction power indicates cases where the model predicts fitness worse than predictions using this average.

The 8-dimensional phenotypic model, which was generated exclusively with the data from subtle environmental perturbations, has substantial predictive power in distant environments (Figure 4). Predictions explain 29% to 95% of the variation in fitness of the 232 test mutants across strong environmental perturbations. For instance, in an environment where glucose concentration was increased from 1.5% to 1.8% and the flask was changed to one that increases the oxygenation of the media (the "Baffle, 1.8% Glucose" environment), we predict 95% of weighted variance with the full 8-component phenotypic model, in contrast to 51% with the 1-component model (Fig 4B). This ability to predict fitness is retained even when the first component (effectively the fitness in EC) is a poor predictor of mutant fitness. For example, in the environment where salt (0.5 M NaCl) was added to the media, the 1-component model predicts fitness worse than predictions based on the average fitness for this environment, resulting in negative variance explained (Fig 4A and 4B). This is due to the fact that mutant fitness in this environment reflects extensive genotype-by-environment interactions, such that the fitness of mutants in this environment is uncorrelated with EC fitness. However, our predictions of mutant fitness in the 0.5 M NaCl environment improve when made using the 8-component phenotypic model, which predicts 72% of weighted variance. Astoundingly, the 8-component model captures strong tradeoffs between mutants with high fitness in the EC and very low fitness in this high salt environment, specifically for *IRA1-nonsense* and, to a lesser extent, *PDE2* mutants (Fig 4B). This was surprising because there appears to be very little variation in fitness of these mutants across the subtle compared to the strong perturbations (Fig 2C).

This ability to predict fitness is also observed for mutations in genes and pathways that are not represented in the 60 that comprise the training set (*e.g.* those with mutations in TOR/Sch9 and HOG pathway genes). For example, the 8-component model explains 93% of variation in the "Baffle, 1.8% Glucose" environment and 71% of variation in the 0.5M NaCl environment for these mutations, compared to 76% and 31% variance explained for the 1-component model, respectively. This indicates that our model is able to capture shared phenotypic effects that extend beyond gene identity. Altogether, our ability to accurately predict the fitness of new mutants in new environments suggests that the phenotypes our model identifies reflect causal effects on fitness.

Most strikingly, phenotypic models that include the three smallest phenotypic components, which together contribute only 1.1% to variance explained across the subtle environmental perturbations (Fig 4A), often explain a substantial amount of variance in the distant environments (Fig 4A; lower panel). For example, the three minor components contribute 17% of the overall weighted variance explained in the 1-Day condition ($\tilde{R}^2$ = 0.6 - 5-component model, $\tilde{R}^2$ = 0.73 - 8-component model; (0.73-0.6)/0.73 = 0.17) and 45% in the 6-Day environment, ($\tilde{R}^2$ = 0.25 - 5-component model, $\tilde{R}^2$ = 0.46 - 8-component model) (Fig 4A and 4B). In contrast, for other strong environments (e.g. Baffle - 1.8% Glucose, 8.5uM GdA (B9) and Baffle - 2.5% Glucose), the three smallest components do not add much explanatory power (Fig 4A). These observations demonstrate that phenotypic components that make very small contributions to fitness in the EC can contribute substantially to fitness in other environments. Overall, these observations suggest an answer to questions about how adaptation is possible when mutations have collateral effects on multiple phenotypes: not all of those phenotypes contribute substantially to fitness in the EC (Fig 1C).

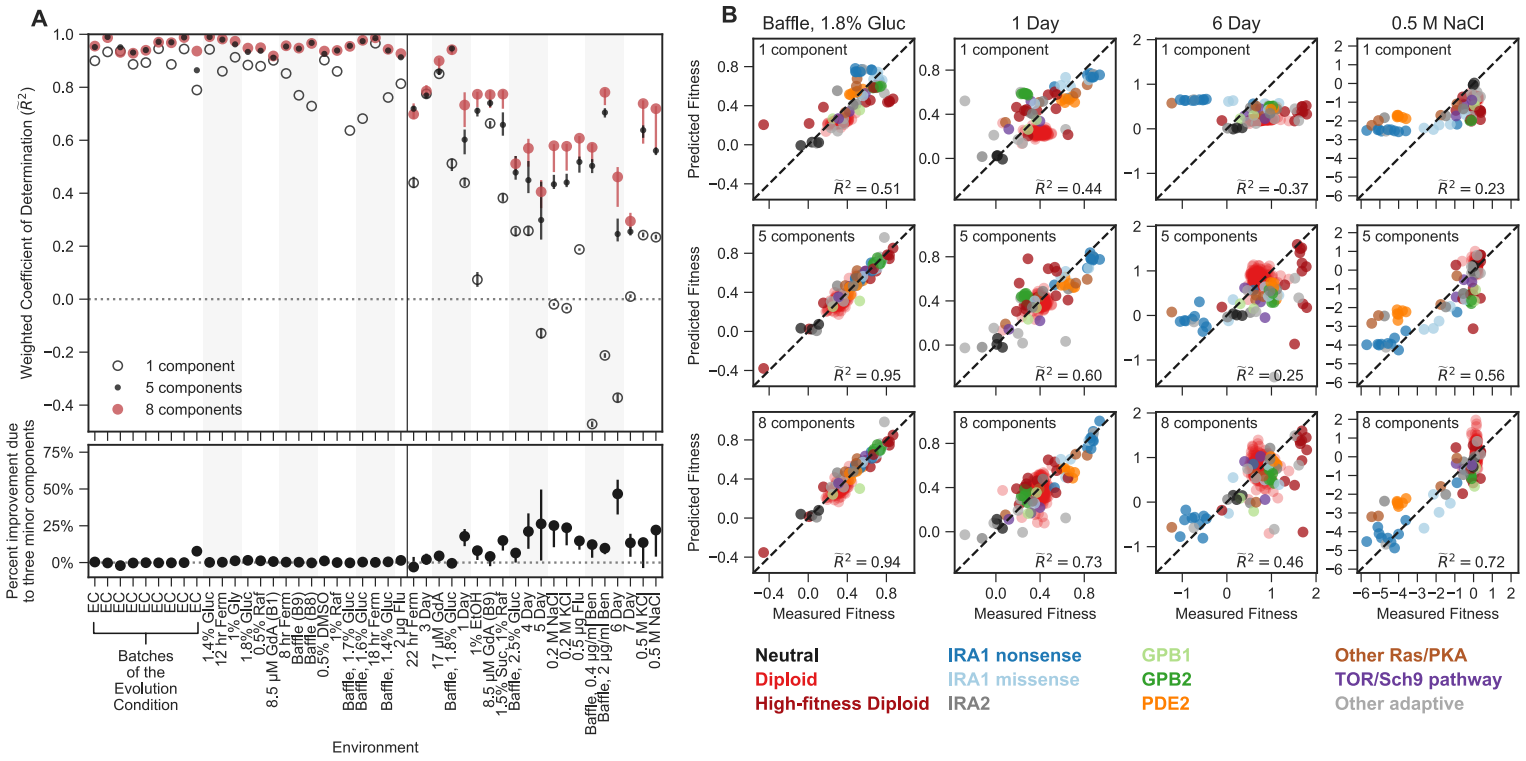

Figure 4. Mutant fitness variation across subtly different environments predicts mutant fitness in novel and substantially different environments. (A) Top panel vertical axis shows the accuracy of fitness predictions in each of 45 environments on the horizontal axis. The accuracy is calculated as the coefficient of determination, weighted such that each mutation type contributes equally. The left side of this plot represents predictions of mutant fitness in subtle environmental perturbations. These predictions are generated by holding out data from that environment when building the phenotypic model. The right side of the plot displays predictions of mutant fitness in strong environmental perturbations. These predictions are generated using a phenotypic model inferred from fitness variation across all 25 subtle different environments (denoted by each of the points or open circles) and for each of the 25 leave-one-out models (range of predictions is depicted with the error bars surrounding each point or open circle). Predictions from the 8-component model (red point) are typically better than the 1-component mode (open circle) and sometimes better than the 5-component model (black point). **Bottom panel** vertical axis shows the percent of the 8-component model's improvement due to the three minor components (calculated by the percent difference between the 5- and 8- component models). The left side shows the improvement of the prediction in subtle environmental perturbations when that subtle perturbation was held out. The right side shows the improvement of the prediction in strong environmental perturbations when using the full model (dots) or the 25 leave-one-out models (the error bars represent the range of improvement). **(B)** For each subplot, the horizontal axis shows the measured fitness value. The vertical axis shows the predicted fitness value when predictions are made using the 1-component (top row), 5-component (middle row), or 8-component (bottom row) models. Columns represent different environments. Points are colored by the mutation type. Note that $\widetilde{R}^2$ less than zero indicates that the prediction is worse than predictions using the mean fitness in that condition (see Methods).

Not all mutants affect all phenotypes and not all environments make all phenotypes important

Next we explore the extent to which the contribution of a phenotypic component to fitness is isolated to a specific environment and/or a specific type of mutation (Fig 5). We find that many phenotypic components matter more to fitness in some environments than others. For instance, component 2 adds on average 36% of the weighted variance in fitness across strong perturbations, despite adding only 7% on average across the subtle environmental perturbations. This contribution is, however, variable, with the second component adding over 90% of variance explained for the two environments with Benomyl and Baffled flasks (the "Baffle, 0.4 µg/ml Benomyl" and "Baffle, 2 µg/ml Benomyl" environments) and only 0.3% for the environment in which the transfer time was lengthened from two to three days (Fig 5A).

This environment-dependence is also true for the smallest two components. Specifically, predictions of mutant fitness in the 0.5 M NaCl environment are improved from the inclusion of component 7, adding 7.5% to weighted variance explained (Fig 5A). Predictions of mutant fitness in the 6 Day transfer environment show improvement from the inclusion of the 8th component, which adds over 15% to weighted variance explained (Fig 5A). However, the predictions of fitness in the 6 Day environment are not improved from the inclusion of the 7th component and the predictions in 0.5 M NaCl are not improved markedly by the inclusion of the 8th component (Fig 5A). This suggests that the phenotypic effects represented by these small components contribute substantially in some environments and not others.

We further asked whether these effects are not only environment-specific but also mutant-specific. To do so, we focused on environments for which the two smallest components contribute substantially to fitness (*e.g.* 0.5 M NaCl). We looked at the extent to which each of these components improves power to predict the fitness of each of the 232 held-out mutants. We found these components improve the fitness predictions for some classes of mutants far more than for others. For example, fitness predictions for mutations in *GPB2,* diploids with chromosome 11 amplifications, and high-fitness diploids with no known mutations each improved by over 4 standard deviations of measurement error in the 0.5 M NaCl environment due to the inclusion of the 7th component (Fig 5B). This phenotypic component also has importance in the 1-Day transfer environment, albeit to a lesser degree, resulting in improvements of roughly 1 standard deviation for each of these mutation types. This suggests that these mutants have some phenotypic effect that contributes only slightly to fitness in many environments, including those that represent subtle perturbations of the EC, but that are particularly important in the 0.5 M NaCl and 1 Day transfer environments. Similarly, we find that the 8th component also improves predictive power for specific types of mutants in specific environments. In this case, diploids with chromosome 11 amplifications and *PDE2* mutants have particularly strong improvements in the 6-Day transfer environment (11 and 5 standard deviations, respectively) and thus likely have a shared phenotypic effect that is captured by component 8 (Fig 5B).

In sum, not all mutants affect all eight phenotypic components to the same degree and not all phenotypic components contribute substantially to fitness in all environments. This idiosyncrasy suggests that directional selection, *e.g.* selection for improved fitness under glucose limited conditions, can generate rather than reduce phenotypic diversity. Though adaptation winnows phenotypic diversity by choosing mutants that affect similar fitness-relevant phenotypes, these mutants may have latent effects on a larger number of phenotypes. When the environment changes, these latent phenotypic effects may have larger contributions to fitness, thus revealing the phenotypic diversity generated by the adaptive process.

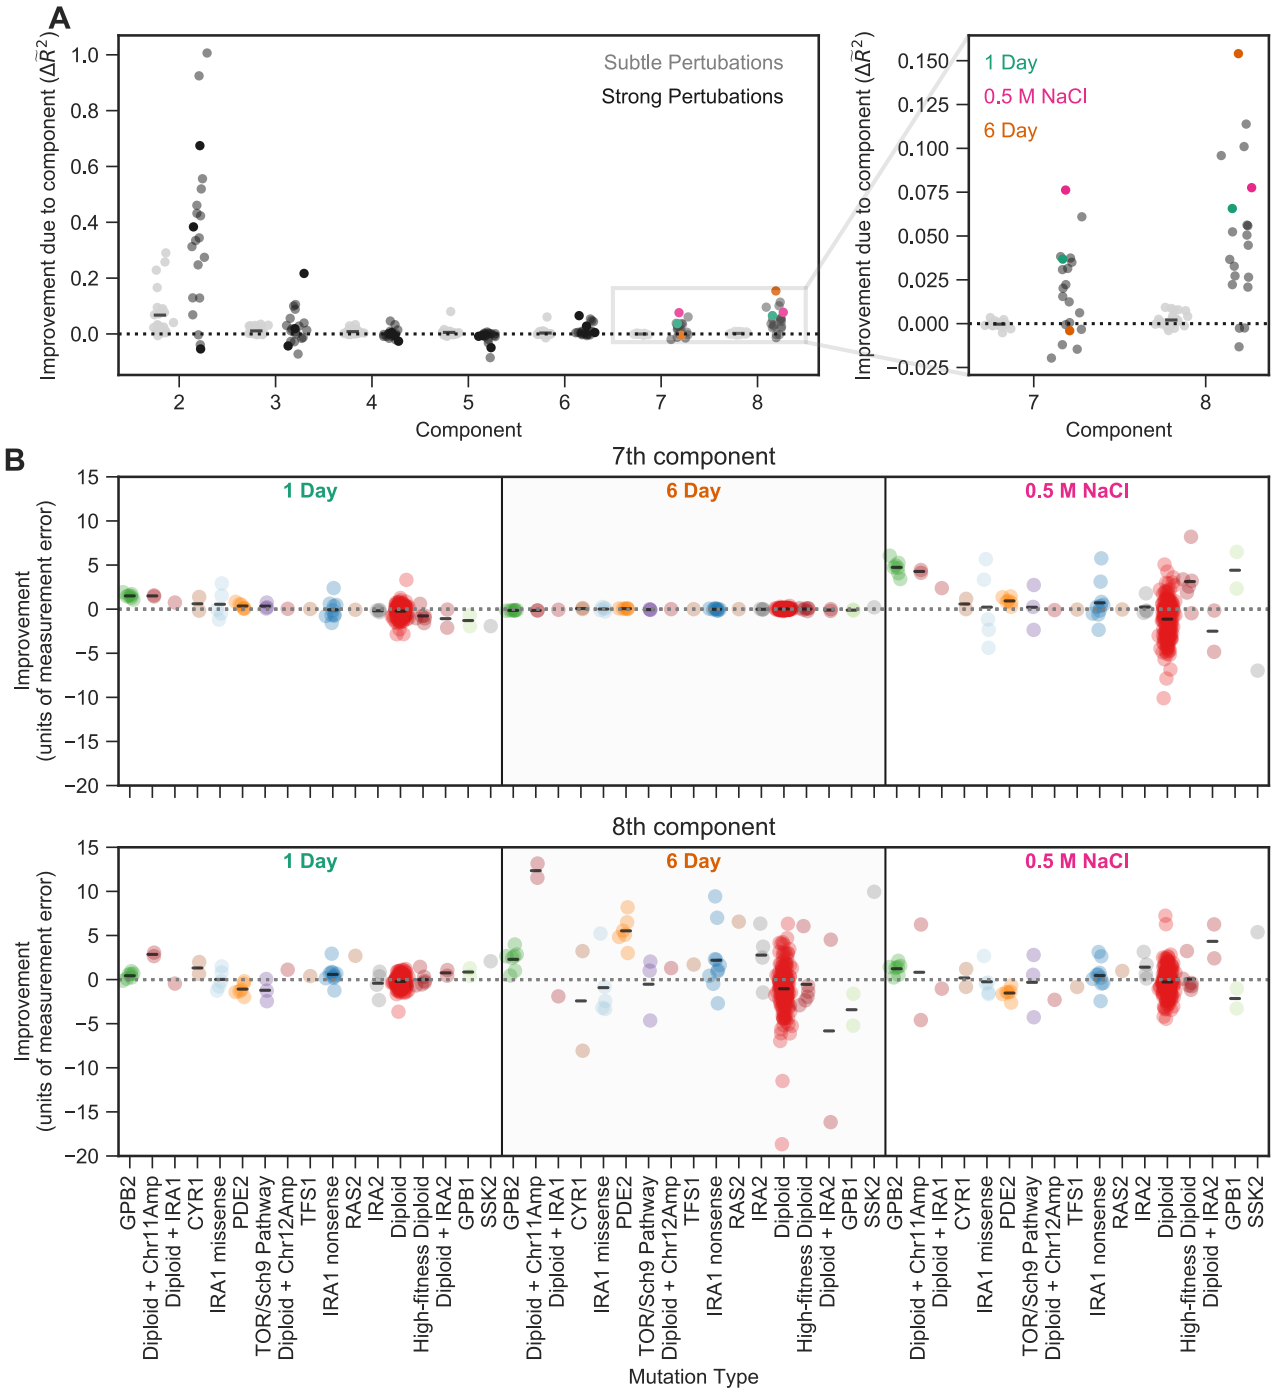

Figure 5. The contribution of a phenotypic component to fitness changes across environments and differs for different types of mutants. (A) Some phenotypic components improve fitness predictions in some environments substantially more than they do in others. The vertical axis shows the improvement in the predictive power of our 8-component phenotypic model due to the inclusion of each component. For example, the improvement due to component 7 is calculated by the difference between the 7-component model and the 6-component model. The improvement of predictive power for each of the subtle environmental perturbations is shown as a gray point and for each of the strong perturbations in black. Magnification shows improvement upon including each of the two smallest components, with three strong perturbations highlighted. **(B)** Some phenotypic components improve fitness predictions for some mutants substantially more than they do for others. For example, the 7th component explains little variation in the 6-Day environment, but the 8th component explains a lot of variation in fitness in the 6-Day environment and is particularly helpful in predicting the fitness of Diploid + Chromosome 11 Amplification mutations in this environment. Vertical axis shows the improvement in predictive power (in units of standard deviation of measurement error) for each type of mutant (denoted on the horizontal axis) in one of three environments (1 Day, 6 Day, and 0.5 M NaCl) when adding either the 7th (top panel) or the 8th (bottom panel) component. Mutants are ordered by the improvement due to the 7th component in the 1 Day environment. Since some types of mutants are more common, e.g. diploids, there are more data points in that category.

DISCUSSION

Here we succeeded in building a low-dimensional statistical model that captures the relationship from genotype to phenotype to fitness for hundreds of adaptive mutants. Mapping the complete phenotypic and fitness impacts of genetic change is a key goal of biology. Such a map is important in order to make meaningful predictions from genetic data (*e.g.* personalized medicine) and to investigate the structure of biological systems (*e.g.* their degree of modularity and pleiotropy) (500 Genomes Field Experiment Team et al., 2019; Collet et al., 2018; Eguchi et al., 2019; Zan and Carlborg, 2020). Our model allows us to do both of these things. We made accurate predictions about the fitness of unstudied mutants across multiple environments, and we gained novel insights about the nature of pleiotropy of the adaptive process. Specifically, we learned that adaptation is modular in the sense that hundreds of diverse adaptive mutants collectively influence a small number of phenotypes that matter to fitness in the evolution condition. We also learned that different mutants have distinct pleiotropic side effects that matter to fitness in other conditions.

Building genotype-phenotype-fitness maps of adaptation has long been an elusive goal due to both conceptual and technical difficulties. Indeed, the very first part of this task, namely the identification of causal adaptive mutations, presents a substantial technical challenge (500 Genomes Field Experiment Team et al., 2019; Barrett et al., 2019, 2008). Fortunately, in some systems such as in microbial experimental evolution and studies of cancer and resistance in microbes and viruses, genomic methodologies combined with availability of repeated evolutionary trials allow us now to detect with high confidence specific genetic changes responsible for adaptation. In the context of microbial evolution experiments, lineage tracing and genomics have opened up the possibility of not only detecting hundreds of specific adaptive events but also measuring their fitness precisely and in bulk (Good et al., 2017; Levy et al., 2015; Li et al., 2019, 2018; Nguyen Ba et al., 2019; Venkataram et al., 2016). Thus in these cases we are coming close to solving the technical challenge of building the genotype to fitness map of adaptation.

However, adding phenotype into this map remains a huge challenge even despite substantial progress in mapping genotype to phenotype (Burga et al., 2019; Camp et al., 2019; Exposito-Alonso et al., 2018; Geiler-Samerotte et al., 2016; Jakobson and Jarosz, 2019; Lee et al., 2019; Paaby et al., 2015; Yengo et al., 2018; Ziv et al., 2017). In principle, we now have advanced tools to measure a large number of phenotypic impacts of a genetic change, for instance through high throughput microscopy, proteomics, or RNAseq (Manzoni et al., 2018; Ritchie et al., 2015; Zhang and Kuster, 2019). The conceptual problem is how to define phenotypes given the interconnectedness of biological systems (Geiler-Samerotte et al., 2019; Paaby and Rockman, 2013). If a mutation leads to complex changes in cell size and shape, should each change be considered a distinct phenotype? Or if a single mutation changes the expression of hundreds or thousands of genes, should we consider each change as a separate phenotype? Intuitively, it seems that we should seek higher order, more meaningful descriptions. For example, perhaps these expression changes are coordinated and reflect the up-regulation of a stress-response pathway. Unfortunately, defining the functional units in which a gene product participates remains difficult, especially because these units re-wire across genetic backgrounds, environments, and species (Geiler-Samerotte et al., 2019; Pavličev et al., 2017; Sun et al., 2020; Zan and Carlborg, 2020).

If mutations influence more than one phenotype, then the mapping from phenotype to fitness also becomes challenging. To investigate this map, we would need to find an artificial way to perturb one phenotype without perturbing others such that we could isolate and measure effects on fitness. Mapping phenotype to fitness is further complicated by the environmental dependence of these relationships (Fragata et al., 2019). For example, a mutation that affects a

cell's ability to store carbohydrates for future use might matter far more in an environment where glucose is re-supplied every 6 days instead of every 48 hours.

In our study, we turned the challenge of environment-dependence into the solution to the seemingly intractable problem of interrogating the phenotype layer of the genotype-phenotype-fitness map. We rely on the observation that the relative fitness of different mutations changes across environments. We assume that differences in how mutant fitness varies across environments must stem from differences in the phenotypes each mutation affects. Rather than *a priori* defining the phenotypes that we think may matter, we use the similarities and dissimilarities in the way fitness of multiple mutants vary across environments to define phenotypes abstractly via their causal effects on fitness. This allows us to dispense with measuring the phenotypes themselves and instead focus on measuring fitness with high precision and throughput, since tools for doing so already exist (Venkataram et al., 2016). This approach has the disadvantage of not identifying phenotypes in a traditional, more transparent way. Still, it represents a major step forward in building genotype-phenotype-fitness maps because it makes accurate predictions and provides novel insights about the phenotypic structure of the adaptive response.

We successfully implemented this approach using a large collection of adaptive mutants evolved in a glucose-limited condition. The first key result is that the map from adaptive mutant to phenotype to fitness is modular, such that it is possible to create a genotype to phenotype to fitness model that is low dimensional. Indeed, our model detects a small number (8) of fitness-relevant phenotypes, the first two of which explain almost all of the variation in fitness (98.3%) across 60 adaptive mutants in 25 environments representing subtle perturbations of the glucose-limited evolution condition. This suggests that the hundreds of adaptive mutations we study — including mutations in multiple genes in the Ras/PKA and TOR/Sch9 pathways, genome duplication (diploidy), and various structural mutations — influence a small number of phenotypes that matter to fitness in the evolution condition. This observation is consistent with theoretical considerations suggesting that mutations that affect a large number of fitness-relevant phenotypes are not likely to be adaptive. It also explains findings from other high-replicate laboratory evolution experiments and studies of cancer that show hundreds of unique adaptive mutations tend to hit the same genes and pathways repeatedly. Our work confirms the intuition that these mutations all affect similar higher-order phenotypes (*e.g.* the level of activity of a signaling pathway). This suggests that, despite the genetic diversity among adaptive mutants, adaptation may be predictable and repeatable at the phenotypic level.

Note that although we detect only 8 fitness-relevant phenotypes, we expect the true number to be much larger as the detectable number is limited by the precision of measurement. We expect this partly because we know that if we had worse precision in this experiment we would have detected fewer than 8 phenotypic components (Fig 3). Still, these additional undetected components cannot be very consequential in terms of their contribution to fitness in the evolution condition, given how well the first 8 components capture variation in environments that are similar to the evolution condition.

Surprisingly, the model built only using subtle environmental perturbations was also predictive of fitness in environments that perturbed fitness strongly. Indeed in some of these environments, such as the environment where 0.5 M NaCl was added to the media or the time of transfer was extended from two to six days, many of the mutants are no longer adaptive and some of them become strongly deleterious. Here the fitness of the mutants in the evolving condition is a very poor predictor of fitness but the full 8-dimensional phenotypic model explains from 29% to 95% of the variance. What was particularly interesting is that the explanatory power of different dimensions was very different for the strong compared to subtle perturbations. For instance, the second dimension which explained 7% of weighted variation on average in the subtle

perturbations, explained 36% on average in the environments that represent strong perturbations. The pattern was particularly striking for the smallest dimensions which at times explained 15% in the strong environmental perturbations while again explaining at most 1% in the subtle environments.

This discovery emphasizes that, although the smaller phenotypic dimensions contribute very little to fitness in the evolution condition (Fig 1B), they can at times have a much larger contribution in other environments (Fig 1C). This makes intuitive sense. For instance, we know that some of the strongest adaptive mutations in our experiment, the nonsense mutations in IRA1, appear to stop cells from shifting their metabolism towards carbohydrate storage when glucose levels become low (Li et al., 2018). This gives these cells a head start once glucose again becomes abundant and does not appear to come at a substantial cost, at least not until these cells are exposed to stressful environments (*e.g.* high salt or long stationary phase) (Li et al., 2018). This example, and more generally the observation that phenotypic effects that are unimportant in the evolving condition can become much more important in other environments, supports the idea that adaptation can happen through large effect mutations because many of the pleiotropic phenotypic effects will be inconsequential in the local environment (Fig 1B – C). We can thus argue that the low-dimensional phenotypic space near the evolution condition hides *latent* and *consequential* phenotypic complexity across the collection of locally phenotypically similar mutants. This complexity is hidden from natural selection in the evolution condition but becomes important once the mutants leave the local environment and are assessed globally for fitness effects. Thus, with respect to their effects on fitness-relevant phenotypes, adaptive mutants may be locally modular, but globally pleiotropic.

The notion of latent phenotypic complexity is exciting as it generates a mechanism by which directional selection generates rather than removes diversity. This suggests a solution to long-standing questions in evolutionary biology about how diversity persists despite directional selection (Walsh and Blows, 2009). Directional selection may promote multiple mutants that affect similar fitness-relevant phenotypes in the evolution condition, but each mutant could have disparate meaningful phenotypic effects that do not contribute immediately to fitness. When the environment changes, these latent phenotypic effects may now matter, allowing for diverse solutions to a variety of possible new environments. This latent phenotypic complexity also has the potential to alter the future adaptive paths that a population takes even in a constant environment. Indeed, these phenotypically diverse mutants are likely to affect the subsequent direction of adaptation given that subsequent mutations can shift the context in which phenotypes are important in the same way as do environmental perturbations. The end result is that directional selection can enhance diversity both within a population and across populations adapting to the same stressors.

The phenomenon of latent phenotypic complexity being driven by adaptation is dependent on there being multiple mutational solutions to an environmental challenge, such that different adaptive mutations might have different latent phenotypic effects. Latent phenotypic diversity might be less apparent in cases where adaptation proceeds through mutations in a single gene and certainly would not exist if adaptation relies on one unique mutation. Thus, in some ways latent phenotypic diversity reflects redundancies in the mechanisms that allow cells to adapt to a challenge. One such putative redundancy in the case investigated in this paper is that the Ras/PKA pathway can be constitutively activated by loss of function mutations to a number of negative regulators including *IRA1*, *PDE2*, and *GPB2*. Mutations in these genes might be redundant in the sense that they influence the same fitness-relevant phenotype in the evolution condition, which in this case is likely flux through the Ras/PKA pathway. This type of redundancy is commonly observed in laboratory evolutions, as evidenced by studies that combine adaptive mutants to find they are no more adaptive than the most fit single mutant (Tenaillon et al., 2012) and the observation that subsequent adaptive mutations tend to be in

other pathways (Aggeli et al., 2020). The major insight from our paper is that we show that mutations with redundant effects on fitness in the evolution condition are not necessarily identical because they may influence different latent phenotypes. This observation adds to a long list of examples demonstrating that redundancies, such as gene duplications and dominance, allow evolution the flexibility to generate diversity.

One disadvantage of our approach is that the phenotypic components that we infer from our fitness measurements are abstract. They represent causal effects on fitness, rather than measurable features of cells. For this reason, perhaps we should not refer to them as phenotypes but rather "fitnotypes" (a mash of the terms "fitness" and "phenotype") that act much like the causal traits in Fisher's geometric model (Fisher, 1930; Harmand et al., 2017; Lourenço et al., 2011; Martin and Lenormand, 2006; Tenaillon, 2014; Tenaillon et al., 2007) or a selectional pleiotropy model (Paaby and Rockman, 2013). Despite this limitation, these fitnotypes have proven useful in allowing us to understand the consequences of adaptive mutation. In addition to insights explained above, we also learned that adaptive mutants in the same gene do not always affect the same fitnotypes. For example, we found that *IRA1-missense* mutations have varied and distinct effects from *IRA1-nonsense* mutations. Further, we believe that identifying fitnotypes will ultimately prove useful in identifying the phenotypic effects of mutation. The fitnotypes can serve as a scaffold onto which a large number of phenotypic measurements can be mapped. Even though fitnotypes are independent with respect to their contribution of fitness, and contribute to fitness linearly, the mapping of commonly measured features of cells (*e.g.* growth rate, the expression levels of growth supporting proteins like ribosomes) onto fitnotypes may not be entirely straightforward. Nonetheless, methods such as Sparse Canonical Correlation Analysis (Suo et al., 2017) hold promise in such a mapping and might help us can relate traditional phenotypes to fitnotypes.

An important question for future research is how common is the pattern we detected in this study and whether it applies to other cases of adaptation in other systems. The method we described is generic and can be applied to any system as long as the fitness of a substantial set of mutants can be profiled for fitness across a moderately large set of environments. This is becoming possible to do in many systems. It is also already clear that the notion that diverse genetic changes can have redundant effects in one environment but distinct and consequential effects in other environments is important to our understanding of adaptation in other settings, including in the context of antibiotic resistance and cancer. Indeed in cancer, tumors even within a particular type of cancer, say lung adenocarcinoma, tend to be extremely genetically diverse even if considering only driver mutations (The Cancer Genome Atlas Research Network, 2014) . The driver mutations do fall into a smaller number of key driver genes and even fewer pathways. While this apparent redundancy might suggest that all tumors are in fact functionally similar in that they solve a small set of functional challenges (Hanahan and Weinberg, 2011, 2000), the notion of latent diversity we propose here suggests that the specific paths taken by the tumors early might matter once the tumor encounters a new challenge such as when they are treated by a cancer therapy. Substantial heterogeneity of tumor response to therapy is consistent with this notion.

Despite the accumulation of large amounts of genomic and phenomic data, integrating this information to identify the phenotypic consequences of mutation that are ultimately responsible for fitness remains incredibly challenging. Our approach allows us to create an abstract representation of the causal effects of genetic mutation and their changing contribution to fitness across environments. This top-down view provides an opening to solving this problem, and combining these approaches with phenotypic measurements will allow us to answer age-old questions about the structure of biological systems and adaptation in a conceptually new and technically powerful and high-throughput way.

ACKNOWLEDGMENTS

The authors thank Sandeep Venkataram for the BarcodeCounter2 script; Yuping Li, Monica Sanchez, Tuya Yokoyama, Chris McFarland, and Dimitra Aggeli for technical assistance; Atish Agarwala, Marc Salit, Sasha Levy, Gavin Sherlock, and all members of the Petrov Lab for helpful comments and discussions. We are grateful to the twitter community that followed #1BigBatch and provided us with very helpful feedback. We are grateful to Enrico Coen for the suggestion of the term “fitnotype”. Some of the computing for this project was performed on the Sherlock cluster. We would like to thank Stanford University and the Stanford Research Computing Center for providing computational resources and support that contributed to these research results. This work was supported by National Institutes of Health grant R35GM118165 (to DAP) and National Institutes of Health grant R35GM133674 (to KGS).

REFERENCES

500 Genomes Field Experiment Team, Exposito-Alonso, M., Burbano, H.A., Bossdorf, O., Nielsen, R., Weigel, D., 2019. Natural selection on the Arabidopsis thaliana genome in present and future climates. Nature 573, 126–129. https://doi.org/10.1038/s41586-019-1520-9

Aggeli, D., Li, Y., Sherlock, G., 2020. Changes in the distribution of fitness effects and adaptive mutational spectra following a single first step towards adaptation (preprint). Evolutionary Biology. https://doi.org/10.1101/2020.06.12.148833

Altenberg, L., 2005. Modularity in Evolution: Some Low-Level Questions, in: Modularity: Understanding the Development and Evolution of Complex Natural Systems. MIT Press, p. 32.

Baeza-Centurion, P., Miñana, B., Schmiedel, J.M., Valcárcel, J., Lehner, B., 2019. Combinatorial Genetics Reveals a Scaling Law for the Effects of Mutations on Splicing. Cell 176, 549-563.e23. https://doi.org/10.1016/j.cell.2018.12.010

Barghi, N., Tobler, R., Nolte, V., Jakšić, A.M., Mallard, F., Otte, K.A., Dolezal, M., Taus, T., Kofler, R., Schlötterer, C., 2019. Genetic redundancy fuels polygenic adaptation in Drosophila. PLOS Biol. 17, e3000128. https://doi.org/10.1371/journal.pbio.3000128

Barrett, R.D.H., Laurent, S., Mallarino, R., Pfeifer, S.P., Xu, C.C.Y., Foll, M., Wakamatsu, K., Duke-Cohan, J.S., Jensen, J.D., Hoekstra, H.E., 2019. Linking a mutation to survival in wild mice. Science 363, 499–504. https://doi.org/10.1126/science.aav3824

Barrett, R.D.H., Rogers, S.M., Schluter, D., 2008. Natural Selection on a Major Armor Gene in Threespine Stickleback. Science 322, 255–257. https://doi.org/10.1126/science.1159978

Boyle, E.A., Li, Y.I., Pritchard, J.K., 2017. An Expanded View of Complex Traits: From Polygenic to Omnigenic. Cell 169, 1177–1186. https://doi.org/10.1016/j.cell.2017.05.038

Burga, A., Ben-David, E., Lemus Vergara, T., Boocock, J., Kruglyak, L., 2019. Fast genetic mapping of complex traits in C. elegans using millions of individuals in bulk. Nat. Commun. 10, 2680. https://doi.org/10.1038/s41467-019-10636-9

Camp, J.G., Platt, R., Treutlein, B., 2019. Mapping human cell phenotypes to genotypes with single-cell genomics. Science 365, 1401–1405. https://doi.org/10.1126/science.aax6648

Chesmore, K., Bartlett, J., Williams, S.M., 2018. The ubiquity of pleiotropy in human disease. Hum. Genet. 137, 39–44. https://doi.org/10.1007/s00439-017-1854-z

Collet, J.M., McGuigan, K., Allen, S.L., Chenoweth, S.F., Blows, M.W., 2018. Mutational Pleiotropy and the Strength of Stabilizing Selection Within and Between Functional Modules of Gene Expression. Genetics 208, 1601–1616. https://doi.org/10.1534/genetics.118.300776

Coombes, D., Moir, J.W.B., Poole, A.M., Cooper, T.F., Dobson, R.C.J., 2019. The fitness challenge of studying molecular adaptation. Biochem. Soc. Trans. 47, 1533–1542. https://doi.org/10.1042/BST20180626

Cowperthwaite, M.C., Bull, J.J., Meyers, L.A., 2005. Distributions of Beneficial Fitness Effects in RNA. Genetics 170, 1449–1457. https://doi.org/10.1534/genetics.104.039248

Crozat, E., Winkworth, C., Gaffe, J., Hallin, P.F., Riley, M.A., Lenski, R.E., Schneider, D., 2010. Parallel Genetic and Phenotypic Evolution of DNA Superhelicity in Experimental Populations of Escherichia coli. Mol. Biol. Evol. 27, 2113–2128. https://doi.org/10.1093/molbev/msq099

Diss, G., Lehner, B., 2018. The genetic landscape of a physical interaction. eLife 7, e32472. https://doi.org/10.7554/eLife.32472

Domingo, J., Baeza-Centurion, P., Lehner, B., 2019. The Causes and Consequences of Genetic Interactions (Epistasis). Annu. Rev. Genomics Hum. Genet. 20, 433–460. https://doi.org/10.1146/annurev-genom-083118-014857

Eckart, C., Young, G., 1936. The approximation of one matrix by another of lower rank. Psychometrika 1, 211–218. https://doi.org/10.1007/BF02288367

Eguchi, Y., Bilolikar, G., Geiler-Samerotte, K., 2019. Why and how to study genetic changes

with context-dependent effects. Curr. Opin. Genet. Dev. 58–59, 95–102. https://doi.org/10.1016/j.gde.2019.08.003
Exposito-Alonso, M., Vasseur, F., Ding, W., Wang, G., Burbano, H.A., Weigel, D., 2018. Genomic basis and evolutionary potential for extreme drought adaptation in Arabidopsis thaliana. Nat. Ecol. Evol. 2, 352–358. https://doi.org/10.1038/s41559-017-0423-0
Fisher, R.A., 1930. The genetical theory of natural selection. Clarendon Press, Oxford. https://doi.org/10.5962/bhl.title.27468
Fragata, I., Blanckaert, A., Dias Louro, M.A., Liberles, D.A., Bank, C., 2019. Evolution in the light of fitness landscape theory. Trends Ecol. Evol. 34, 69–82. https://doi.org/10.1016/j.tree.2018.10.009
Geiler-Samerotte, K.A., Li, S., Lazaris, C., Taylor, A., Ziv, N., Ramjeawan, C., Paaby, A.B., Siegal, M.L., 2019. Extent and context dependence of pleiotropy revealed by high-throughput single-cell phenotyping (preprint). Evolutionary Biology. https://doi.org/10.1101/700716
Geiler-Samerotte, K.A., Zhu, Y.O., Goulet, B.E., Hall, D.W., Siegal, M.L., 2016. Selection Transforms the Landscape of Genetic Variation Interacting with Hsp90. PLOS Biol. 14, e2000465. https://doi.org/10.1371/journal.pbio.2000465
Good, B.H., McDonald, M.J., Barrick, J.E., Lenski, R.E., Desai, M.M., 2017. The dynamics of molecular evolution over 60,000 generations. Nature 551, 45–50. https://doi.org/10.1038/nature24287
Hanahan, D., Weinberg, R.A., 2011. Hallmarks of Cancer: The Next Generation. Cell 144, 646–674. https://doi.org/10.1016/j.cell.2011.02.013
Hanahan, D., Weinberg, R.A., 2000. The Hallmarks of Cancer. Cell 100, 57–70. https://doi.org/10.1016/S0092-8674(00)81683-9
Harmand, N., Gallet, R., Jabbour-Zahab, R., Martin, G., Lenormand, T., 2017. Fisher's geometrical model and the mutational patterns of antibiotic resistance across dose gradients: FISHER'S GEOMETRICAL MODEL AND ANTIBIOTIC RESISTANCE. Evolution 71, 23–37. https://doi.org/10.1111/evo.13111
Illumina, 2017. Effects of Index Misassignment on Multiplexing and Downstream Analysis [WWW Document]. URL https://www.illumina.com/content/dam/illumina-marketing/documents/products/whitepapers/index-hopping-white-paper-770-2017-004.pdf?linkId=36607862. (accessed 6.25.20).
Jakobson, C.M., Jarosz, D.F., 2019. Molecular Origins of Complex Heritability in Natural Genotype-to-Phenotype Relationships. Cell Syst. 8, 363-379.e3. https://doi.org/10.1016/j.cels.2019.04.002
Josse, J., Sardy, S., 2014. Adaptive Shrinkage of singular values. ArXiv13106602 Stat.
Kacser, H., Burns, J.A., 1981. The Molecular Basis of Dominance. Genetics 97, 639–666.
Karageorgi, M., Groen, S.C., Sumbul, F., Pelaez, J.N., Verster, K.I., Aguilar, J.M., Hastings, A.P., Bernstein, S.L., Matsunaga, T., Astourian, M., Guerra, G., Rico, F., Dobler, S., Agrawal, A.A., Whiteman, N.K., 2019. Genome editing retraces the evolution of toxin resistance in the monarch butterfly. Nature 574, 409–412. https://doi.org/10.1038/s41586-019-1610-8
Kemble, H., Eisenhauer, C., Couce, A., Chapron, A., Magnan, M., Gautier, G., Le Nagard, H., Nghe, P., Tenaillon, O., 2020. Flux, toxicity, and expression costs generate complex genetic interactions in a metabolic pathway. Sci. Adv. 6, eabb2236. https://doi.org/10.1126/sciadv.abb2236
Lang, G.I., Rice, D.P., Hickman, M.J., Sodergren, E., Weinstock, G.M., Botstein, D., Desai, M.M., 2013. Pervasive genetic hitchhiking and clonal interference in forty evolving yeast populations. Nature 500, 571–574. https://doi.org/10.1038/nature12344
Langmead, B., Salzberg, S.L., 2012. Fast gapped-read alignment with Bowtie 2. Nat. Methods 9, 357–359. https://doi.org/10.1038/nmeth.1923
Lee, J.T., Coradini, A.L.V., Shen, A., Ehrenreich, I.M., 2019. Layers of Cryptic Genetic Variation Underlie a Yeast Complex Trait. Genetics 211, 1469–1482.

https://doi.org/10.1534/genetics.119.301907
Levy, S.F., Blundell, J.R., Venkataram, S., Petrov, D.A., Fisher, D.S., Sherlock, G., 2015. Quantitative evolutionary dynamics using high-resolution lineage tracking. Nature 519, 181–186. https://doi.org/10.1038/nature14279
Li, C., Zhang, J., 2018. Multi-environment fitness landscapes of a tRNA gene. Nat. Ecol. Evol. 2, 1025–1032. https://doi.org/10.1038/s41559-018-0549-8
Li, Y., Petrov, D.A., Sherlock, G., 2019. Single nucleotide mapping of trait space reveals Pareto fronts that constrain adaptation. Nat. Ecol. Evol. 3, 1539–1551. https://doi.org/10.1038/s41559-019-0993-0
Li, Y., Venkataram, S., Agarwala, A., Dunn, B., Petrov, D.A., Sherlock, G., Fisher, D.S., 2018. Hidden Complexity of Yeast Adaptation under Simple Evolutionary Conditions. Curr. Biol. 28, 515-525.e6. https://doi.org/10.1016/j.cub.2018.01.009
Lind, P.A., Farr, A.D., Rainey, P.B., 2015. Experimental evolution reveals hidden diversity in evolutionary pathways. eLife 4, e07074. https://doi.org/10.7554/eLife.07074
Lourenço, J., Galtier, N., Glémin, S., 2011. COMPLEXITY, PLEIOTROPY, AND THE FITNESS EFFECT OF MUTATIONS. Evolution 65, 1559–1571.
Manzoni, C., Kia, D.A., Vandrovcova, J., Hardy, J., Wood, N.W., Lewis, P.A., Ferrari, R., 2018. Genome, transcriptome and proteome: the rise of omics data and their integration in biomedical sciences. Brief. Bioinform. 19, 286–302. https://doi.org/10.1093/bib/bbw114
Martin, G., Lenormand, T., 2006. A GENERAL MULTIVARIATE EXTENSION OF FISHER'S GEOMETRICAL MODEL AND THE DISTRIBUTION OF MUTATION FITNESS EFFECTS ACROSS SPECIES. Evolution 60, 893–907.
McInnes, L., Healy, J., Melville, J., 2018. UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction. ArXiv180203426 Cs Stat.
Mehlhoff, J.D., Stearns, F.W., Rohm, D., Wang, B., Tsou, E.-Y., Dutta, N., Hsiao, M.-H., Gonzalez, C.E., Rubin, A.F., Ostermeier, M., 2020. Collateral fitness effects of mutations. Proc. Natl. Acad. Sci. 117, 11597–11607. https://doi.org/10.1073/pnas.1918680117
Melo, D., Porto, A., Cheverud, J.M., Marroig, G., 2016. Modularity: Genes, Development, and Evolution. Annu. Rev. Ecol. Evol. Syst. 47, 463–486. https://doi.org/10.1146/annurev-ecolsys-121415-032409
Nguyen Ba, A.N., Cvijović, I., Rojas Echenique, J.I., Lawrence, K.R., Rego-Costa, A., Liu, X., Levy, S.F., Desai, M.M., 2019. High-resolution lineage tracking reveals travelling wave of adaptation in laboratory yeast. Nature 575, 494–499. https://doi.org/10.1038/s41586-019-1749-3
Orr, H.A., 2000. Adaptation and The Cost of Complexity. Evolution 54, 13–20. https://doi.org/10.1111/j.0014-3820.2000.tb00002.x
Otwinowski, J., McCandlish, D.M., Plotkin, J.B., 2018. Inferring the shape of global epistasis. Proc. Natl. Acad. Sci. 115, E7550–E7558. https://doi.org/10.1073/pnas.1804015115
Owen, A.B., Perry, P.O., 2009. Bi-cross-validation of the SVD and the nonnegative matrix factorization. Ann. Appl. Stat. 3, 564–594. https://doi.org/10.1214/08-AOAS227
Paaby, A.B., Rockman, M.V., 2013. The many faces of pleiotropy. Trends Genet. 29, 66–73. https://doi.org/10.1016/j.tig.2012.10.010
Paaby, A.B., White, A.G., Riccardi, D.D., Gunsalus, K.C., Piano, F., Rockman, M.V., 2015. Wild worm embryogenesis harbors ubiquitous polygenic modifier variation. eLife 4, e09178. https://doi.org/10.7554/eLife.09178
Pavličev, M., Wagner, G.P., Chavan, A.R., Owens, K., Maziarz, J., Dunn-Fletcher, C., Kallapur, S.G., Muglia, L., Jones, H., 2017. Single-cell transcriptomics of the human placenta: inferring the cell communication network of the maternal-fetal interface. Genome Res. 27, 349–361. https://doi.org/10.1101/gr.207597.116
Pressman, A.D., Liu, Z., Janzen, E., Blanco, C., Müller, U.F., Joyce, G.F., Pascal, R., Chen, I.A., 2019. Mapping a Systematic Ribozyme Fitness Landscape Reveals a Frustrated Evolutionary Network for Self-Aminoacylating RNA. J. Am. Chem. Soc. 141, 6213–6223.

https://doi.org/10.1021/jacs.8b13298
Ritchie, M.D., Holzinger, E.R., Li, R., Pendergrass, S.A., Kim, D., 2015. Methods of integrating data to uncover genotype–phenotype interactions. Nat. Rev. Genet. 16, 85–97. https://doi.org/10.1038/nrg3868
Rockman, M.V., 2012. THE QTN PROGRAM AND THE ALLELES THAT MATTER FOR EVOLUTION: ALL THAT'S GOLD DOES NOT GLITTER. Evolution 66, 1–17. https://doi.org/10.1111/j.1558-5646.2011.01486.x
Sarkisyan, K.S., Bolotin, D.A., Meer, M.V., Usmanova, D.R., Mishin, A.S., Sharonov, G.V., Ivankov, D.N., Bozhanova, N.G., Baranov, M.S., Soylemez, O., Bogatyreva, N.S., Vlasov, P.K., Egorov, E.S., Logacheva, M.D., Kondrashov, A.S., Chudakov, D.M., Putintseva, E.V., Mamedov, I.Z., Tawfik, D.S., Lukyanov, K.A., Kondrashov, F.A., 2016. Local fitness landscape of the green fluorescent protein. Nature 533, 397–401. https://doi.org/10.1038/nature17995
Sella, G., Barton, N.H., 2019. Thinking About the Evolution of Complex Traits in the Era of Genome-Wide Association Studies. Annu. Rev. Genomics Hum. Genet. 20, 461–493. https://doi.org/10.1146/annurev-genom-083115-022316
Sinha, R., Stanley, G., Gulati, G.S., Ezran, C., Travaglini, K.J., Wei, E., Chan, C.K.F., Nabhan, A.N., Su, T., Morganti, R.M., Conley, S.D., Chaib, H., Red-Horse, K., Longaker, M.T., Snyder, M.P., Krasnow, M.A., Weissman, I.L., 2017. Index switching causes "spreading-of-signal" among multiplexed samples in Illumina HiSeq 4000 DNA sequencing (preprint). Molecular Biology. https://doi.org/10.1101/125724
Sivakumaran, S., Agakov, F., Theodoratou, E., Prendergast, J.G., Zgaga, L., Manolio, T., Rudan, I., McKeigue, P., Wilson, J.F., Campbell, H., 2011. Abundant Pleiotropy in Human Complex Diseases and Traits. Am. J. Hum. Genet. 89, 607–618. https://doi.org/10.1016/j.ajhg.2011.10.004
Starr, T.N., Flynn, J.M., Mishra, P., Bolon, D.N.A., Thornton, J.W., 2018. Pervasive contingency and entrenchment in a billion years of Hsp90 evolution. Proc. Natl. Acad. Sci. 115, 4453–4458. https://doi.org/10.1073/pnas.1718133115
Sun, S., Baryshnikova, A., Brandt, N., Gresham, D., 2020. Genetic interaction profiles of regulatory kinases differ between environmental conditions and cellular states. Mol. Syst. Biol. 16. https://doi.org/10.15252/msb.20199167
Suo, X., Minden, V., Nelson, B., Tibshirani, R., Saunders, M., 2017. Sparse canonical correlation analysis. ArXiv170510865 Stat.
Tenaillon, O., 2014. The Utility of Fisher's Geometric Model in Evolutionary Genetics. Annu. Rev. Ecol. Evol. Syst. 45, 179–201. https://doi.org/10.1146/annurev-ecolsys-120213-091846
Tenaillon, O., Rodriguez-Verdugo, A., Gaut, R.L., McDonald, P., Bennett, A.F., Long, A.D., Gaut, B.S., 2012. The Molecular Diversity of Adaptive Convergence. Science 335, 457–461. https://doi.org/10.1126/science.1212986
Tenaillon, O., Silander, O.K., Uzan, J.-P., Chao, L., 2007. Quantifying Organismal Complexity using a Population Genetic Approach. PLoS ONE 2, e217. https://doi.org/10.1371/journal.pone.0000217
The Cancer Genome Atlas Research Network, 2014. Comprehensive molecular profiling of lung adenocarcinoma. Nature 511, 543–550. https://doi.org/10.1038/nature13385
Venkataram, S., Dunn, B., Li, Y., Agarwala, A., Chang, J., Ebel, E.R., Geiler-Samerotte, K., Hérissant, L., Blundell, J.R., Levy, S.F., Fisher, D.S., Sherlock, G., Petrov, D.A., 2016. Development of a Comprehensive Genotype-to-Fitness Map of Adaptation-Driving Mutations in Yeast. Cell 166, 1585-1596.e22. https://doi.org/10.1016/j.cell.2016.08.002
Verduyn, C., Postma, E., Scheffers, W.A., Van Dijken, J.P., 1992. Effect of benzoic acid on metabolic fluxes in yeasts: A continuous-culture study on the regulation of respiration and alcoholic fermentation. Yeast 8, 501–517. https://doi.org/10.1002/yea.320080703
Visscher, P.M., Yang, J., 2016. A plethora of pleiotropy across complex traits. Nat. Genet. 48, 707–708. https://doi.org/10.1038/ng.3604

Wagner, G.P., Altenberg, L., 1996. Perspective: Complex Adaptations and the Evolution of Evolvability. Evolution 50, 967–976. https://doi.org/10.2307/2410639
Wagner, G.P., Pavlicev, M., Cheverud, J.M., 2007. The road to modularity. Nat. Rev. Genet. 8, 921–931. https://doi.org/10.1038/nrg2267
Wagner, G.P., Zhang, J., 2011. The pleiotropic structure of the genotype–phenotype map: the evolvability of complex organisms. Nat. Rev. Genet. 12, 204–213. https://doi.org/10.1038/nrg2949
Walsh, B., Blows, M.W., 2009. Abundant Genetic Variation + Strong Selection = Multivariate Genetic Constraints: A Geometric View of Adaptation. Annu. Rev. Ecol. Evol. Syst. 40, 41–59. https://doi.org/10.1146/annurev.ecolsys.110308.120232
Weinreich, D.M., 2006. Darwinian Evolution Can Follow Only Very Few Mutational Paths to Fitter Proteins. Science 312, 111–114. https://doi.org/10.1126/science.1123539
Welch, J.J., Waxman, D., 2003. MODULARITY AND THE COST OF COMPLEXITY. Evolution 57, 1723–1734. https://doi.org/10.1111/j.0014-3820.2003.tb00581.x
Yengo, L., Sidorenko, J., Kemper, K.E., Zheng, Z., Wood, A.R., Weedon, M.N., Frayling, T.M., Hirschhorn, J., Yang, J., Visscher, P.M., the GIANT Consortium, 2018. Meta-analysis of genome-wide association studies for height and body mass index in ~700000 individuals of European ancestry. Hum. Mol. Genet. 27, 3641–3649. https://doi.org/10.1093/hmg/ddy271
Zan, Y., Carlborg, Ö., 2020. Dynamic genetic architecture of yeast response to environmental perturbation shed light on origin of cryptic genetic variation. PLOS Genet. 16, e1008801. https://doi.org/10.1371/journal.pgen.1008801
Zhang, B., Kuster, B., 2019. Proteomics Is Not an Island: Multi-omics Integration Is the Key to Understanding Biological Systems. Mol. Cell. Proteomics 18, S1–S4. https://doi.org/10.1074/mcp.E119.001693
Ziv, N., Shuster, B.M., Siegal, M.L., Gresham, D., 2017. Resolving the Complex Genetic Basis of Phenotypic Variation and Variability of Cellular Growth. Genetics 206, 1645–1657. https://doi.org/10.1534/genetics.116.195180

METHODS

LEAD CONTACT AND MATERIALS AVAILABILITY

Further information and requests for resources and reagents should be directed to and will be fulfilled by the Lead Contact, Dmitri Petrov (dpetrov@stanford.edu).

EXPERIMENTAL MODEL AND SUBJECT DETAILS

The yeast strains used in this study can be grown and maintained using standard methods (e.g. YPD media in test tubes, glycerol stocks for long term storage at -80°C), but should be propagated in the appropriate selection environment (a glucose-limited minimal media - M3 medium for the evolution condition) for comparable fitness and phenotypic measurements. All of the strains we study are of genetic background MATα, ura3Δ0, ybr209w::Gal-Cre-KanMX-1/2URA3-loxP-Barcode-1/2URA3-HygMX-lox66/71.

Experiments were performed with barcoded mutants isolated from a previous evolution experiment (Levy et al., 2015). To measure their fitness, these mutants were competed against a constructed reference strain with a restriction site in the barcode region (Venkataram et al., 2016).

Since we utilize some data from previous experiments (Li et al., 2018; Venkataram et al., 2016), the collection of adaptive barcoded mutants that we studied differed slightly across environments. These differences can be thought of as another parameter that varies across the environments (*e.g.* in addition to glucose or salt concentration). In some experiments, we used a collection containing 4,800 adaptive mutants that do not necessarily start at equal frequency (Venkataram et al., 2016). In other experiments, we used a collection containing a subset of 500 of these mutants where each one starts at equal frequency (Li et al., 2018; Venkataram et al., 2016). Though the smallest collection of mutants we study comprises 500 strains, our work focuses on 292 of these (Table S1). We focus on strains for which we obtained fitness measurements in 45 environments and for which mutations conferring fitness advantages have been previously identified, either by whole genome sequencing or using a drug to test ploidy (Li et al., 2018; Venkataram et al., 2016).

In a few experiments, we spiked in re-barcoded mutants and additional neutral lineages as internal controls. Since re-barcoded mutants are identical, except for the barcode, these teach us about the precision with which we can measure a mutant's fitness. Specifically, we spiked in ten re-barcoded *IRA1-nonsense* mutants (each with a frameshift insertion AT to ATT mutation at bp 4090) and ten *IRA1-missense* mutants (each with a G to T mutation at bp 3776). Neutral lineages teach us about the behavior of the unmutated reference strain, which we must infer because it's barcode is eliminated from the experiment before sequencing. The spiked in neutrals include ten barcoded lineages from the original evolution experiment (Levy et al., 2015) for which whole genome sequencing did not reveal any mutations (Venkataram et al., 2016) and previous fitness measurements did not reveal any deviation from the reference (Li et al., 2018; Venkataram et al., 2016).

METHOD DETAILS

Conducting the barcoded fitness measurements

Fitness measurement experiments were performed as described previously (Li et al., 2018; Venkataram et al., 2016), where growth competitions were set up between a pool of barcoded mutants and a reference strain. The change in the frequency of each barcode over time reflects the fitness of the adaptive mutant possessing that barcode, relative to the reference strain.

We conducted fitness measurements under a variety of conditions (Table S2) that represent perturbations of the condition in which these adaptive mutants evolved. Briefly, we separately grew up an overnight culture of the barcode pool and the ancestral reference strain in 100mL M3 (minimal, glucose-limited) medium (Verduyn et al., 1992). We then mixed these saturated cultures at a 1:9 ratio such that 90% of cells represent the reference strain. This ratio allows for mutants to compete against the ancestor rather than competing against each other. We then inoculated 400μL of this mixed culture ($\sim 5 \times 10^7$ cells) into 100mL of fresh media in 500mL DeLong flasks. The type of media used, and sometimes the shape of the flask, varied depending on condition (Table S2). This culture was then grown at 30°C in an incubator shaking at 223 RPM for 48 hours. After 48 hours of growth, 400μL of saturated culture was transferred into fresh media of the same type, in a new flask of the same type. This serial dilution was usually continued 4 times, yielding 5 time-points over which to measure the rate at which a barcode's frequency changed, though some experiments include one more or one less depending on the experimenter and on whether technical problems (*e.g.* PCR failure) caused loss of time-points.

After each transfer of 400 μL, the left-over 9600μL was frozen so that we could later sequence the barcodes present at every time-point. To prepare this culture for freezing, it was transferred to 50mL conicals, spun down at 3000 rpm for 5 minutes, resuspended in 5mL of sorbitol freezing solution (0.9M sorbitol, 0.1M Tris-HCL pH 7.5, 0.1M EDTA pH 8.0), aliquoted into three 1.5mL tubes, and stored at -80°C.

For experiments where additional neutral lineages and re-barcoded lineages were included, the initial inoculation mix consisted of 90% ancestral reference strain, 9.4% barcode mutant pool, 0.2% additional neutral spike-in pool, 0.2% re-barcoded IRA1 nonsense pool, and 0.2% re-barcoded IRA1 missense pool.

Growth conditions

In this study, we present fitness measurement data from a collection of 45 conditions that each represent perturbations of the growth condition in which these adaptive mutants evolved. We refer to this original evolution condition as the "EC". In the EC, cells are grown in flasks with a flat bottom and transferred to new flasks every 48 hours (see *Conducting the barcoded fitness measurements*). Cells are grown in M3 media (Verduyn et al., 1992). This media is glucose-limited, meaning the cells run out of glucose before any other nutrient. In the EC, the starting glucose concentration is 1.5%.

The 45 perturbations of the EC are summarized in Table S2 and include changes to the growth media, the flask shape, and the transfer times. For example, in the "1 Day" condition, we change the transfer time from 48 to 24 hours. In the "1.8% glucose, baffled flask condition" we change the starting glucose concentration from 1.5% to 1.8% and change the flask type from one with a flat bottom to one with baffles. Several of these conditions include experiments from previous studies (Li et al., 2018; Venkataram et al., 2016).

For each of these 45 conditions but three, we include between two and four replicates that were performed simultaneously (Table S2) such that overall we performed a total of 109 fitness measurements on our collection of adaptive mutants. Our replicate structure is nested in that some of our 45 conditions represent replicate experiments that we performed at different times. Variation across experiments performed at different times is often referred to as "batch effects" and likely reflects environmental variability that we were unable to control (e.g. slight fluctuations in incubation temperature due to limits on the precision of the instrument). In particular, we re-measured the fitness of the adaptive mutants in the EC on 9 different occasions, each time including 3 or more replicates. We refer to these 9 experiments as 'EC batches' in the main text. However, every set of experiments that was performed at the same time constitutes a separate "batch". There were slight differences across batches in the way we prepared barcodes for sequencing, which we detail in the relevant Methods sections. This variation across batches can be thought of as another parameter that varies across the 45 conditions (in addition to glucose or salt concentration). We report which experiments were performed in the same batch in Table S2.

Some conditions, including some Fluconazole conditions and Geldanamycin conditions, have unexpected orderings in the strength of perturbation (i.e. the smaller drug concentration shows a larger difference in fitness or similar concentrations seem to have different effects). Regardless of whether these observations reflect technical problems (*e.g.* degradation or poor solubility of the drug), we include these conditions because we use the effect of the realized perturbation on fitness to build low-dimensional phenotypic models. In other words, the identity of the perturbation does not matter in this study.

DNA Extraction of each sample

After a growth competition is complete, we extracted DNA from frozen samples following either a protocol described previously (for batches 1 – 6 and 10) (Venkataram et al., 2016) or a modified protocol that improves the ease and yield of extraction. Our modified protocol is as follows. For each sample, a single tube of the three that were frozen for each sample (see *Conducting the barcoded fitness measurements*) was removed from the freezer and thawed at room temperature. We extracted DNA from that sample using the following modification of the Lucigen MasterPure yeast DNA purification kit (#MPY80200). We transferred the thawed cells into a 15mL conical and centrifuge for 3 min at 4000 RPM. After discarding the supernatant, the pellet was then resuspended with 1.8 mL of the MasterPure lysis buffer, and 0.5 mm glass beads were added to help with disruption of the yeast cell wall. The mix of pellet, lysis buffer, and beads was then vortexed for 10 seconds and incubated for 45 minutes at 65°C, with periodic vortexing. The solution was then put on ice for 5 min and then 900 μL of MPC Protein Reagent was mixed with the solution. We then separated protein and cell debris by centrifugation at 4000 RPM, transferring 1900μL of supernatant to a 2 mL centrifuge tube. We further separated remaining protein and cell debris by centrifuging at 13200 RPM for 5 min. The supernatant was then divided into two 2mL centrifuge tubes, with 925 μL of the supernatant into each. Next, we added 1000μL of isopropanol to each tube, mixed by inversion, centrifuged at 13200 RPM for 5 min, and discarded the supernatant. The pellet, containing the DNA was then resuspended in 250μL of Elution Buffer and 10 μL of 5 ng/μL RNAase A was added. This was either left at room temperature overnight or incubated at 60°C for 15 min. Next the two tubes per sample were combined into a single tube and 1500 μL of ethanol was added. This was then mixed by inversion, and strands of precipitating DNA appeared. This was centrifuged at 13200 RPM for 2 min, and the supernatant was discarded. We again precipitated the DNA by resuspending with 750 μL of ethanol, and collected the DNA by centrifuging 13200 RPM for 2 min. The supernatant was discarded, and the tubes were left to air dry. Finally, we resuspended the pellet in Elution Buffer to a final concentration of 50 ng/ul for later use in PCR reactions (approximately 3600 ng of DNA were used for the PCR reactions).

PCR Amplification of the Barcode Locus

After extracting DNA, we PCR-amplified the barcode locus for each sample. Batches 1 – 6 and 10 were conducted with the protocols described in (Li et al., 2018; Venkataram et al., 2016). We made some slight modifications to this protocol, including using a new set of primers to allow for nested-unique-dual index labeling, for batches 7, 8, and 9. Our modified protocol is as follows.

We used a two-step PCR protocol to amplify the barcodes from the DNA. The first PCR cycle uses primers with "inline indices" to label samples (see *Mitigating the effects of index hopping* section for details). These inline indices are highlighted in bold below. Attaching unique indices to samples pertaining to different conditions or timepoints allows us to multiplex these samples on the same sequencing lane. Each primer also contains a Unique Molecular Identifier (UMI) – denoted by the sequence of "N" nucleotides in the primer – which is used to determine if identical barcode sequences each represent yeast cells that were present at the time the sample was frozen, or a PCR amplification of the a barcode from a single cell (see Levy et al., 2015; Li et al., 2018; Venkataram et al., 2016). Primers were HPLC purified to ensure they are the correct length.

Forward primers

Primer Name	Sequence
F201	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **CGATGTT** TAATATGGACTAAAGGAGGCTTTT
F202	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **ACAGTGT** TAATATGGACTAAAGGAGGCTTTT
F203	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **TGACCAT** TAATATGGACTAAAGGAGGCTTTT
F204	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **GCCAATT** TAATATGGACTAAAGGAGGCTTTT
F205	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **ATCACGT** TAATATGGACTAAAGGAGGCTTTT
F206	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **CAGATCT** TAATATGGACTAAAGGAGGCTTTT
F207	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **GGCTACT** TAATATGGACTAAAGGAGGCTTTT
F208	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **TAGCTTT** TAATATGGACTAAAGGAGGCTTTT
F209	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **TTAGGCT** TAATATGGACTAAAGGAGGCTTTT
F210	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **ACTTGAT** TAATATGGACTAAAGGAGGCTTTT
F211	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **GATCAGT** TAATATGGACTAAAGGAGGCTTTT
F212	TCGTCGGCAGCGTC AGATGTGTATAAGAGACAG NNNNNNNN **CTTGTAT** TAATATGGACTAAAGGAGGCTTTT

Reverse primers

Primer Name	Sequence
R301	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **TATATACGC** TCGAATTCAAGCTTAGATCTGATA
R302	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **CGCTCTATC** TCGAATTCAAGCTTAGATCTGATA
R303	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **GAGACGTCT** TCGAATTCAAGCTTAGATCTGATA
R304	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **ATACTGCGT** TCGAATTCAAGCTTAGATCTGATA
R305	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **ACTAGCAGA** TCGAATTCAAGCTTAGATCTGATA
R306	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **TGAGCTAGC** TCGAATTCAAGCTTAGATCTGATA
R307	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **CTGCTACTC** TCGAATTCAAGCTTAGATCTGATA
R308	GTCTCGTGGGCTCGG AGATGTGTATAAGAGACAG NNNNNNNN **GCGTACGCA** TCGAATTCAAGCTTAGATCTGATA

For the first step of PCR, we performed 8 reactions per sample to offset the effects of PCR jackpotting within each reaction. For each set of 8 reactions, we used the master mix:

- 200μL OneTaq Hot Start 2X Master Mix with Standard Buffer (NEB M0484L)
- 8μL 10uM Forward primer
- 8μL 10uM Reverse primer
- 72μL sample genomic DNA (diluted to 50ng/μL or all of sample if between 25-50ng/μL)
- 16μL 50mM MgCl2
- 96μL Nuclease Free Water (Fisher Scientific #AM9937)

We then aliquoted 50μL of the master mix into each of 8 PCR tubes, and ran on the thermocycler with the following cycle:

1. 94°C for 10 min
2. 94°C for 3 min
3. 55°C for 1 min
4. 68°C for 1 min
5. Repeat steps 2-4 2x (for a total of 3 cycles)
6. 68°C for 1 min
7. Hold at 4°C

We then added 100μL of binding buffer from the ThermoScientific GeneJET Gel Extraction Kit (#K0692) to each PCR reaction, and performed a standard PCR purification protocol in one column per sample. In the final step, we eluted into 80μL of elution buffer.

For the second step of PCR, we use standard Nextera XT Index v2 primers (Illumina #FC-131-2004) to further label samples representing different conditions and timepoints with unique identifiers that allow for multiplexing on the same sequencing lane. We uniquely dual-indexed each sample using our nested scheme (see *Mitigating the effects of index hopping* section for details). We performed 3 reactions of the second step PCR per sample, using the master mix:

- 1.5μL Q5 Polymerase (NEB #M0491L)
- 30μL Q5 Buffer (NEB #M0491L)
- 3μL 10mM dNTP (Fisher Scientific #PR-U1515)
- 6.25μL i7 Nextera XT Primer (“N” primer)
- 6.25μL i5 Nextera XT Primer (“S” primer)
- 78μL purified step 1 PCR product
- 25μL Nuclease Free Water (Fisher Scientific #AM9937)

This master mix was then divided into 3 PCR tubes per reaction, and run with the following protocol on a thermocycler:

1. 98°C for 30 sec
2. 98°C for 10 sec
3. 62°C for 20 sec
4. 72°C for 30 sec
5. Repeat steps 2-4 at least 21 times and at most 27 times (for a total of 22 to 28 cycles)
6. 72°C for 3 min
7. Hold at 4°C

We then added 100μL of binding buffer from the ThermoScientific GeneJET Gel Extraction Kit and purified the PCR product, eluting into 43μL. We found that increasing the number of cycles in the second step PCR beyond 21 did not seem to improve the amount of DNA recovered after

gel extraction. For some samples, we experimented with a touch down procedure for the second step PCR where we started with a hotter annealing temperature and slowly decreased it over the course of 27 cycles. This also did not seem to increase the yield of DNA recovered from the PCR.

Removal of the Reference Strain via Digestion and Gel Purification

To avoid the vast majority of our sequencing reads mapping only to the reference strain (and thus not being informative to relative fitness of the mutants), we use restriction digest to cut the ApaLI restriction site in the middle of the reference strain's barcode region. We mixed 43μL of the second step PCR product with 2μL of ApaLI (NEB #R0507L) and 5μL of 10X Cutsmart and incubated at 37°C for at least 2 hours (up to overnight). After digestion, we conducted size selection by running the digested sample on a gel, removing all product less than 300bp, and isolating the DNA using a standard ThermoScientific GeneJET Gel Extraction protocol. Our expected product is 350bp. We did not remove longer sequences via gel extraction because of the possibility that some barcode sequences may selectively form complexes with themselves or other barcodes.

Note that for some samples, we also digested the reference strain before PCR, in addition to after PCR, to decrease the amount of reference strain barcode. For these samples, we mixed 80ul of genomic DNA (at concentration 50ng/μL) with 10μL of 10X Cutsmart and 2μL of ApaLI and incubated 37°C for at least 2 hours (up to overnight). This product was then used as the template for PCR step 1 (with appropriate water volume adjustments to ensure 50μL reactions).

Sample pooling and Amplicon Sequencing

We used the Qubit High Sensitivity (ThermoFisher #Q32854) method to quantify the concentration of the final product for each sample, then pooled samples with different dual indices in equal frequency for sequencing. Our samples were then sent to either Novogene (https://en.novogene.com/) or Admera Health (https://www.admerahealth.com/) for quality control (qPCR and either Bioanalyzer or TapeStation) and sequencing. We used 2x150 paired-end sequencing along with index sequencing reads on Illumina HiSeq machines using patterned flow cells (either HiSeq 4000 or HiSeq X). We also used Illumina Nextseq machines with unpatterned flow cells. We found that the former was more subject to index hopping errors, please see *Mitigating the effects of index hopping* for a discussion of how our dual indexing reduces effects of index hopping. All amplicon samples were sequenced with at least 20% genomic DNA spiked in (either whole genomes from an unrelated project or phi-X) to ensure adequate diversity on the flow cell.

Mitigating the effects of index hopping

To reduce the effects of index hopping observed on Illumina patterned flow cell technology (including HiSeq 4000, HiSeq X, and Novaseq machines) (Illumina, 2017; Sinha et al., 2017), we devise a nested unique-dual-indexing approach. This approach uses a combination of inline indices attached during the first step of PCR, as well as Nextera indices attached during the second step of PCR. The latter indices are not part of the sequencing read (they are read in a separate Index Read). This process uniquely labels both ends of all DNA strands such that DNA strands from multiple samples can be multiplexed on the same flow cell. Had we only labeled one end of each DNA strand, index hopping could have caused us to incorrectly identify some reads as coming from the wrong sample.

One approach to label samples with unique-dual-indices is to use 96 forward primers, each of which is paired to one of 96 reverse primers, instead our nested approach allows us to uniquely dual-index samples with only 40 total primers (12 forward inline, 8 reverse inline, 12 Nextera i7, 8 Nextera i5). Specifically, we can use combinations of the Nextera and inline primers. One way

to think of this is that there are 96 possible ways to combine the forward inline and Nextera i5 primers that are on the same side of the read, effectively creating 96 unique labels for that end of the read.

To reduce the effect of index hopping contamination on our results, we included only samples that were sequenced on non-patterned flow cell technology (HiSeq 2000 and 2500 for samples in batches 1-6, 10, NextSeq for samples in batch 9) or were sequenced on patterned flow cell technology (patterned flow cell HiSeq) with nested unique-dual indexing.

Processing of Amplicon Sequencing Data

We processed the amplicon sequencing data by first using the index tags to de-multiplex reads representing different conditions and timepoints. Then, using Bowtie2 (Langmead and Salzberg, 2012), we mapped reads to a known list of barcodes generated by Venkataram et al. (2016), removed PCR duplicates using the UMIs from the first-step primers, and counted the number of reads for each barcode in each sample. The source code for this step can be found at https://github.com/sandeepvenkataram/BarcodeCounter2. We processed all raw data for this study using this pipeline, including re-processing the raw sequencing files for data from previous studies (Li et al., 2018; Venkataram et al., 2016) so that all data was processed together using the most recent version of the code.

Several samples included technical replicates where the sample was split at various times in the process, including before DNA extraction, before PCR, and prior to sequencing. Read counts across these technical replicates were merged in order to calculate the best estimate of barcode frequencies. Counts were merged after appropriately accounting for PCR duplicates as identified from Unique Molecular Identifiers.

QUANTIFICATION AND STATISTICAL ANALYSIS

Fitness Estimate Inference

The amplicon sequencing data shows the relative frequency of each barcode in each time-point of every one of our 109 fitness measurement experiments. To estimate the fitness of each barcoded mutant in each experiment, we calculate how barcode frequencies change over time. We do this using previously described methods (Venkataram et al., 2016).

Briefly, we first calculated the log-frequency change of each barcoded adaptive mutant for each subsequent pair of time-points. This log-frequency change must be corrected by the mean fitness of the population, such that it represents the relative fitness of each mutant relative to the reference strain, which makes up the bulk of the population. Since we destroyed barcodes pertaining to the reference strain by digesting them, we infer how the mean fitness of the population changes at each time-point using barcoded lineages that are known to be neutral (see *Identification of neutral lineages*). Once we calculated the change in the relative fitness of each barcoded mutant across each pair of consecutive time-points, we took a weighted average across all pairs as our final estimate of each adaptive mutant's relative fitness for a given experiment. We weighted each pairwise fitness estimate using an uncertainty measure generated from a noise model (see *Noise model* section below).

This results in 109 fitness measurements per each barcoded mutant, with some of the 45 conditions having more representation than others due to having more replicates. In cases where we have replicates, we averaged the fitness values across the replicates, weighted by the measurement uncertainty, resulting in our final 45 fitness estimates per each adaptive mutant lineage.

We included only timepoints with at least 1,000 reads for which at least 400 mutants were detected to have at least 1 read. Furthermore, we required that fitness measurement experiments must include at least three timepoints to be included in our analysis.

Identification of Neutral Lineages

Previous work using this fitness measurement method focused on a larger collection of 4800 barcoded yeast lineages, where the vast majority of these lineages were neutral (Li et al., 2018; Venkataram et al., 2016). In order to increase the number of reads per adaptive lineage, we used a smaller pool of 500 lineages for most experiments. However, this prevents us from identifying neutral lineages as was done in previous studies, by rejecting outlier lineages with higher than typical fitness values. Instead, we used a set of 35 high-confidence neutral lineages to infer mean fitness (see *Experimental model and subject details*). These lineages showed no fitness differences from the neutral expectation in previous studies and were shown to possess no mutations detectable via whole genome sequencing. These high-confidence neutral lineages were present in all experiments, and were spiked into experiments from batch 9 to increase their frequency. We used these neutrals to perform the fitness inference in two steps. First, we inferred fitness using this collection of high-confidence neutrals to make a first pass at inferring the fitness values. Next, we included lineages with similar behavior to the high-confidence neutrals to improve our estimate of mean fitness.

Noise model

To quantify the uncertainty for each fitness measurement, we used the noise model as outlined in Venkataram et al., 2016.

Briefly, this noise model accounts for the uncertainty coming from several sources of noise. The first type of noise scales with the number of reads for a given lineage. This noise stems from stochasticity in population dynamics (coming from the inherent stochasticity in growth and noise associated with dilution), from counting noise associated with a finite coverage, and technical noise from DNA extraction and PCR. We fit this noise by quantifying the variation in the frequency of neutral lineages (see *Identification of neutral lineages*). There is additional variation in fitness observed for high-frequency lineages between replicate experiments (here we refer to variation across replicates that were performed simultaneously, not variation across batches). We also accounted for this uncertainty following previous studies. Specifically, we fit an additional frequency-independent source of noise using between-replicate variation.

Checks on noise model

Because our ability to count the phenotypes that matter to fitness hinges upon measurement error, we further assessed the accuracy of our noise model. We did so by using barcoded lineages that should have the same fitness because they are genetically identical. Since our fitness estimates are imperfect (*i.e.* they contain some noise), we estimated each of these lineages as having slightly different fitness. We then asked if the variation in fitness across these lineages is explained by our noise model, or if there is more variation than our noise model can account for. We did this explicitly by calculating, for each lineage, how far our fitness estimate is from the best guess for the true underlying fitness value (the group's mean) in units of the estimate's measurement precision. We then calculated the percent of lineages that are a given distance from the group's average to understand the accuracy of the model. For instance, if the noise model perfectly captures the uncertainty of each measurement, then 10% of the diploid lineages should have a difference from the weighted diploid mean in the 10th percentile, 20% in the 20th percentile, etc. Because 188 of our 292 barcode mutants are diploids without additional mutations, diploids are an ideal group to use to assess the accuracy of the noise model. This procedure shows that, for the vast majority of replicates, the noise model is conservative. That is, diploid lineages tend to have less variation in fitness than expected by the

noise model (Fig S1).

Classifying mutants by mutation type

Some types of mutants are present more than others. For example, 188 of our 292 mutants are diploids and 30 mutants are in the IRA1 gene. If not properly accounted for, this imbalance can lead to some unfairness in predictions for our model. For example, if we use mostly diploid lineages to train our model, we will be very good at predicting the fitness of diploids but poor at predicting other types of mutants. This means that we must classify our mutants by mutation type in order to properly balance them. We classified mutants following previous work (Venkataram et al., 2016) that classified mutants as either diploids, or if haploid, by the gene possessing the putative causal mutation. Because previous work finds differences in fitness between missense and nonsense/frameshift/indel mutations in IRA1, here we classified these mutants into "missense" and "nonsense" classes, where mutants with frameshift and indel mutations were classified as "nonsense". We also classified diploid mutants with additional mutations in nutrient-response genes or chromosomal amplifications as separate groups. Additionally, we created a separate class for "high-fitness diploid" mutants that possess no additional detected mutations (other than being diploid) but have very high fitness in the EC. To be classified as a high-fitness diploid, a diploid mutant must have an average fitness across all 9 EC batches that is greater than 2 standard deviations above the average of all diploids.

Calculation of Weighted Average Z Score

To partition environments into subtle and strong perturbations of the EC, we relied on the 9 experiments performed in the EC. Since each of these experiments was performed at a different time, variation in fitness across these experiments represents batch effects, and we therefore refer to these 9 experiments as "EC batches". Environmental differences between batches are as subtle as possible, as they represent the limit of our ability to minimize environmental variation. Thus, variation in fitness across the EC batches serves as a natural benchmark for the strength of environmental perturbations. If the deviations in fitness caused by an environmental perturbation are substantially stronger than those observed across the EC batches, we call that perturbation "strong".

More explicitly, to determine whether a given environmental perturbation is subtle or strong, we first quantified the typical variation in fitness for each mutant, across the EC batches:

$$\sigma_i = \frac{1}{n_{batches}} \sum_{j}^{batches} \left| f_{ij} - \overline{f_i} \right|$$

where σ_i^2 represents the variance in fitness across the EC batches for mutant *i*, and $\overline{f_i}$ represents the average fitness of mutant *i* across the EC batches.

To ensure that each mutation type contributes equally to our classification of how different each environment is from the evolution condition, we weighed each mutant's contribution to this difference. We did so based on the number of mutants with the same mutation type, such that the mutation-type-weighted average Z-score for a given environment *j* is given by:

$$z_j = \sum_{i}^{mutants} \frac{\left| f_{ij} - \overline{f_i} \right|}{n_{type(i)} \sigma_i}$$

where $n_{type(i)}$represents the number of mutants that are the same mutation type as mutant *i.*

We then classified the environmental perturbations based on this Z-score. Sixteen environments provoked fitness differences resulting in a Z-score of less than two, and we classified these environmental perturbations as "subtle". The remaining 20 environments had Z-scores greater

than 2, which we classified as "strong" environmental perturbations.

Model of phenotypes that contribute to fitness

In order to count the phenotypes that affect fitness in our collection of mutants, we explored a low-dimensional phenotypic model. We explicitly used a model of fitness-relevant phenotypes such that each mutant is represented as having a fixed effect on each phenotype, represented by a vector of *k* phenotypes, e.g. mutant *i* is represented by the vector $(p_{i1}, p_{i2}, p_{i3}, \ldots, p_{ik})$. In addition, each environment is represented by a vector of phenotypic weights, representing the importance of each of the *k* phenotypes to fitness in that environment, e.g. environment *j* represented by the column vector $(e_{1j}, e_{2j}, e_{3j}, \ldots, e_{kj})$. The fitness effect of mutant *i* in a given environment *j* is the linear combination of that mutant's phenotypes, each weighted by its importance in environment *j*:

$$f_{ij} = p_{i1}e_{1j} + p_{i2}e_{2j} + p_{i3}e_{3j} + \ldots + p_{ik}e_{kj}$$

Our fitness measurements reflect mutant fitness relative to a reference strain, therefore, our model places the reference strain (which has fitness 0 by definition) at the origin of this multi-dimensional space. Our model only includes phenotypes that differ between the reference strain and least one mutant. This is sensible given that our reference strain is a modified version of the ancestor of all of these mutant lineages. Thus, if there exists a phenotype that contributes to fitness, but none of the adaptive mutants altered that phenotype, our model will not detect it. More explicitly, a phenotype that contributes to fitness would have a non-zero value of *e*, but if no mutant alters that phenotype from the reference, all mutants would have a zero value of *p* for that phenotype. Thus, the non-zero value of *e* would always be multiplied by a zero value for *p* and this phenotypic dimension would not be represented in our model. This is not to say that if only a single mutant of the 292 alters a particular phenotype we would include it as a phenotypic dimension. Our power to add dimensions to our model is limited by measurement noise. We only include dimensions that capture more variation in fitness than do dimensions that capture measurement noise (see *Estimating the detection threshold using measurement error*).

Similarly, because we measure fitness, and not phenotype, our model is blind to any phenotypic effect that does not contribute to fitness in at least one of the 45 environments we studied. If a mutant has large phenotypic effects, but they do not cause that mutant's fitness to differ from the reference strain in any of these 45 environments, this phenotypic effect will not be represented in our low-dimensional phenotypic model. More explicitly, mutants may have non-zero phenotypic effects *p*, but if these do not influence their fitness in any environment we study, *e* will be zero for all 45 environments. Thus, *p* times *e* will also be zero and we will not include this phenotypic dimension in our model.

Importantly, the phenotypic dimensions that we infer from our fitness measurements are abstract entities. They represent causal effects on fitness, rather than measurable features of cells. For this reason, they might be called "fitnotypes" (a mash of the terms "fitness" and "phenotype"). Even though the fitnotypes are independent with respect to their contribution of fitness, and contribute to fitness linearly, the mapping of commonly measured features of cells (*e.g.* growth rate, the expression levels of growth supporting proteins like ribosomes) onto fitnotypes may be more complicated. For instance, a commonly measured cellular feature that has a complicated nonlinear mapping to fitness could be detected as many, linearly-contributing fitnotypes. This is another reason that our phenotypic dimensions are not necessarily comparable to what people traditionally think of as a "phenotype".

Using Singular Value Decomposition to decompose the fitness matrix

Our goal is to use fitness measurements to learn about the phenotypic effects of mutations as well as the contribution of these phenotypes to fitness in different environments. We conducted fitness measurements for 292 mutants in each of 45 environments and organized these data

into a fitness matrix, *F*, where every row corresponds to a mutant, every column corresponds to an environment, and every entry is a fitness measurement. Because our model (see *Model of phenotypes that contribute to fitness*) represents fitness in a given environment as the sum of multiple phenotypes, each scaled by their contribution to fitness in that environment, we can use Singular Value Decomposition (SVD) to decompose the fitness matrix Fas:

$$P\Sigma E^T = F$$

The left hand side of this equation consists of three matrices:P, which represents the positions of the mutants in our low-dimensional model of phenotypic space, E^T, which represents the contribution of a phenotype to fitness in a given environment, and Σ, a diagonal matrix representing the singular values of the fitness matrix F. Though the singular values are informative in this separation of three matrices, particularly for the amount of variation captured by each of the inferred components, we can also think of this as a decomposition into two matrices, where we fold the singular values into either the mutant phenotypes or the environment weights, as described in the main text. Either way, this decomposition captures the data represented in the fitness matrix F, including measurement error as well as the underlying biological signals.

Importantly, the dimensions in the phenotypic model we built using SVD are detected in the order of their explanatory power. Moreover, the first dimension is the best, linear 1-component model that explains the data (if evaluated by mean squared error). This is true for any set of the first *k* components. This means, for example, that the model with the first eight components is the best possible 8-component linear model for the observed data (Eckart and Young, 1936).

One issue in this type of analysis is that adding more components always improves the explanatory power of the model, even when those components capture variation that is primarily due to measurement noise. This type of overfitting problem is common in statistics, and several methods have been devised to select the appropriate number of components to include. We use two such methods here.

Estimating the detection threshold using measurement error

One method to select the appropriate number of components to include in the model and prevent overfitting (*i.e.* prevent fitting a component that primarily represents noise) is to use measurement error as a type of control. This is only possible if the amount of measurement error is known. We estimated the amount of noise in our fitness measurements using a previously described noise model (see *Noise Model*) (Venkataram et al., 2016). Since this noise model includes counting noise, every fitness measurement may have a different amount of noise. For example, mutants present at low frequency will be subject to more stochasticity resulting from counting noise. We used this noise model to simulate fitness tables (*F)* where mutant fitnesses vary exclusively due to measurement noise. We simulated 1000 noise-only matrices, where each entry is pulled from a normal distribution centered at zero and with variance equal to the estimated measurement noise of the corresponding entry in the true fitness matrix *F*. We then applied SVD to each noise-only matrix, which gave us a set of singular values generated only by noise. From many such simulations, we took the average size of the largest component, which reveals how much variation can be explained by a component that captures only noise. We found that the largest noise-components are of the size that they would capture 0.07% of variation in our true fitness matrix. Thus, we set this as our limit of detection. In other words, in order for us to include 8 components in our low-dimensional model, all of them must explain more than 0.07%% of the variation in fitness. This approach is analogous to identifying a threshold when measurement noise is known but not identical for all entries in the matrix (Josse and Sardy, 2014).

Estimating detection threshold using bi-cross-validation

Another method for identifying the appropriate number of components is to use their predictive power. This method relies on the intuition that measurement error is uncorrelated across different mutants and different environments. Therefore, a component that represents measurement error should not contain information that can help predict the fitnesses of these mutants in new environments. It should also not contain information that can help predict the fitness of unstudied mutants. We used a bi-cross-validation scheme of the SVD devised by Owen and Perry (2009) which divides the mutants and environments into distinct groups of training and testing sets. This subsequently divided our matrix of fitness measurements into 4 submatrices: the fitness of the training mutants in the training environments (*D)*, the fitness of the training mutants in the testing environments (*C)*, the fitness of the testing mutants in the training environments (*B),* and the fitness of the testing mutants in the testing environments (*A*).

$$F = \begin{pmatrix} A \begin{matrix} Test\ Mutants \\ Test\ Environments \end{matrix} & B \begin{matrix} Test\ Mutants \\ Train\ Environments \end{matrix} \\ C \begin{matrix} Train\ Mutants \\ Test\ Environments \end{matrix} & D \begin{matrix} Train\ Mutants \\ Train\ Environments \end{matrix} \end{pmatrix}$$

We carried out SVD on the training data (submatrix *D*), which returned a set of singular values and corresponding components that captured the fitness data in *D*. We then used these components to predict the fitness of the testing mutants in the testing environments (submatrix *A*). First, we tried to predict these fitness values by only using the first component. That is, we fixed this first component and the first singular value for the training mutants. We then found the best first component for the testing environments based on the fitness values of the training mutants in these environments (i.e. using the information in submatrix *D*), given the constraint that the training mutants can only be represented by the one component. We then conducted an analogous procedure to find the first component of the testing mutants by fixing the first component of the training environments by using the information in submatrix *B*. Then, we tried to predict the fitness of the testing mutants in the testing environments using the first component independently fit for each. We subsequently repeated this procedure, giving the testing mutants access to more of the training components each time. If the components detected by the training components represent biological signal, then this should improve the ability to predict the fitness of the testing mutants in the testing environments. However, once the components primarily represent measurement error, their inclusion should harm predictive power. Therefore, we use the number of components with the best ability to predict the held-out data (submatrix *A)* as the number of components that represent biological signal in our data.

For computational efficiency, we explicitly used the formulation proposed by Owen and Perry (2009) for the prediction of the held-out submatrix *A:*

$$\hat{A} = B\left(\widehat{D}^{(k)}\right)^{+} C$$

where $\left(\widehat{D}^{(k)}\right)^{+}$ denotes the Moore-Penrose inverse of the rank *k* approximation of sub-matrix *D*. This prediction is equivalent to the procedure outlined above, provided that least-squares regression is used to identify the components of the testing mutants and testing conditions, conditional upon the training components (Owen and Perry, 2009).

We divided our mutants into fixed training and testing sets (see *Division of Mutants into Training and Testing Sets*) and used these sets throughout our study. As for training versus testing environments, these changed depending on our goal. For validating the number of components to include in our phenotypic model, we held out each of the 25 subtle environmental perturbations, using it as the testing environment and the other 24 for training. For making predictions of the fitness of the testing mutants in the strong environmental perturbations, we used all 25 subtle environmental perturbations as the training set, though we also show how

these predictions vary when each of the 25 subtle environmental perturbations is held out from the training set.

Division of Mutants into Training and Testing Sets

In order to perform bi-cross-validation on our data, we need to divide our data into training and testing sets. Because some mutation types, in particular diploids and Ras/PKA mutants, are present more than others in our collection of mutants, we sampled the training set such that each mutation type is represented roughly equally (see *Classifying mutants by mutation type*). Specifically, we designated half of each mutation type, with a maximum of 20 representatives of each type, as belonging to the training set. The remaining mutants comprise the test set. For example, there are 188 diploids included in the 292 adaptive mutants. We included 20 in the training set and 168 in the test set. There are 20 *IRA1-nonsense* mutants included in the 292, and we included 10 in the training and 10 in the test set. Additionally, genes that are represented only once in the set of mutations are placed in the test set. This results in a training set of 60 mutants and a testing set of 232 mutants (see Table S1).

Clustering mutants in phenotype space

After inferring the low-dimensional model of phenotype space using SVD, we used Uniform Manifold Approximation and Projection (UMAP) to visualize how the mutants cluster in that space. For this analysis, we used the 8-component phenotypic model that we built from the 60 training mutants and the 25 subtle perturbations. We did this to avoid the model being dominated by variation in very common mutations, specifically the diploids, which make up 188/292 of our adaptive mutants. We added more mutants in the visualization by finding the location of each of the testing mutants (except diploids) by least sum of squares optimization. To do so we fixed the coordinates for the 25 environments and found the coordinates for each mutant that best estimated its fitness in all environments. To further avoid our visualization being dominated by the diploids, we included only the diploids present in the training set in our visualization. For UMAP, we specified that 20 neighbors are used.

Though UMAP tends to preserve both local and global structure (McInnes et al., 2018) it is not necessarily representative of the distance between objects in high-dimensional space. Thus, to quantify more precisely the clustering by gene observed, we explicitly compared the median pairwise distance between these apparent clusters to 10000 randomly chosen sets of the same size and calculated empirical p-values. Because there are many diploids such that they will be the most prevalent type of mutant drawn in these randomly chosen sets, we only drew from strains that have other mutations besides or in addition to diploidy.

Calculation of Weighted Coefficient of Determination

Because mutants are present in unequal numbers in the test set, standard measures of variance explained are likely to be representative of our ability to predict mutants that have many barcoded lineages present in the data, for instance diploid and IRA1-nonsense mutations. These measures would be less representative of mutants with few lineages present, i.e. TOR/Sch9 pathway mutants. Thus, we use a measure of predictability ($\tilde{R}^2$) that weights the contribution of each mutant to overall variance explained based on the number of lineages that share its mutation type (diploids, IRA1 nonsense, IRA1 missense, GPB2, etc.). This effectively measures our ability to predict the fitness of each mutation type, rather than each mutant. For overall predictive power across all mutants and conditions, we used the measure:

$$\tilde{R}^2 = 1 - \frac{\sum_i^{mutants} \sum_j^{conditions} \frac{1}{n_{type(i)}} \left(f_{ij} - \widehat{f_{ij}}\right)^2}{\sum_i^{mutants} \sum_j^{conditions} \frac{1}{n_{type(i)}} \left(f_{ij} - \bar{f}\right)^2}$$

where $\bar{f}$ denotes the average fitness for all evaluated mutants and evaluated conditions.

We used a similar measure to quantify the ability to predict fitness for each environment *j*. This is given by:

$$\widetilde{R_j}^2 = 1 - \frac{\sum_i^{mutants} \frac{1}{n_{type(i)}} \left(f_{ij} - \widehat{f_{ij}}\right)^2}{\sum_i^{mutants} \frac{1}{n_{type(i)}} \left(f_{ij} - \overline{f_j}\right)^2}$$

where $\overline{f_j}$ denotes the average fitness across all evaluated mutants in condition *j.*

Note that this measure explicitly compares a model's fitness prediction in each environment to predictions made using the average fitness in that environment, such that if the model's fitness prediction is the same as the average fitness, $\tilde{R}^2$ is zero. It is possible that a given model's fitness prediction is worse than that of the average fitness in that environment, resulting in negative values of $\tilde{R}^2$. In our work, negative $\tilde{R}^2$ values occur for the 1-component model when predicting the fitness of mutants in some of the strong environmental perturbations. In particular, this occurs when fitness in that environment is uncorrelated with EC fitness, which is captured by the first component, such that the EC fitness is unable to make reasonable predictions of fitness in this environment.

Note that we observe qualitatively similar results to this measure when we use a standard variance explained measure and exclude diploids, which dominate the test set (see Fig S5).

Calculating mutant-specific improvement

It is possible that all 292 of our adaptive mutants each affect all 8 of the phenotypic components in our low-dimensional model, however, it is also possible that some mutants influence some phenotypes more strongly than others. In order to quantify how much a specific component lends to the ability to predict the fitness of each mutant in each environment, we need a metric to calculate the difference in predictive accuracy for the model with and without this component. Specifically, to assess the impact of the inclusion of the *k*th component, we compared the prediction accuracy of the *k*-component model to the model that includes the first *k-1* components.

Because fitness estimates vary in their reliability due to finite coverage and other sources (see *Noise model* section), we should factor this uncertainty in our measure of prediction improvement. For example, a small improvement in prediction accuracy for a very uncertain fitness estimate is less meaningful than the same improvement in prediction accuracy for a fitness estimate that we are quite confident in. Thus, we scale the difference in prediction accuracy by the amount of uncertainty in the underlying fitness estimate.

This gives us the measure of improvement in the estimate of the fitness of mutant *i* in condition *j* due to the inclusion of the *n*th component as:

$$I_{ij}^k = \frac{\left(\widehat{f_{ij}}^{k-1} - f_{ij}\right) - \left(\widehat{f_{ij}}^{k} - f_{ij}\right)}{\epsilon_{ij}}$$

where $\widehat{f_{ij}}^{k}$ and $\widehat{f_{ij}}^{k-1}$ represent the estimate of the fitness of mutant *i* in condition *j* for the model with *k* and *k-1* components, respectively. f_{ij} and ϵ_{ij} represent the measured fitness value and measurement uncertainty for the fitness of mutant *i* in condition *j*, respectively.

DATA AND CODE AVAILABILITY

Data Resource

The raw Illumina sequencing data for the fitness measurement assays conducted in this study can be found under NIH BioProject: PRJNA641718. Sequencing data previously published in Venkataram et al., 2016 can be found under NIH BioProject: PRJNA310010. Sequencing data previously published in Li et al., 2018 can be found under NIH BioProject: PRJNA388215.

Code

The software repository for the barcode counting code can be found at https://github.com/sandeepvenkataram/BarcodeCounter2.
The software repository for the fitness estimate inference can be found at https://github.com/barcoding-bfa/fitness-assay-python.
The code for all downstream analysis, including figure generation can be found at https://github.com/grantkinsler/1BigBatch.

SUPPLEMENTARY FIGURES

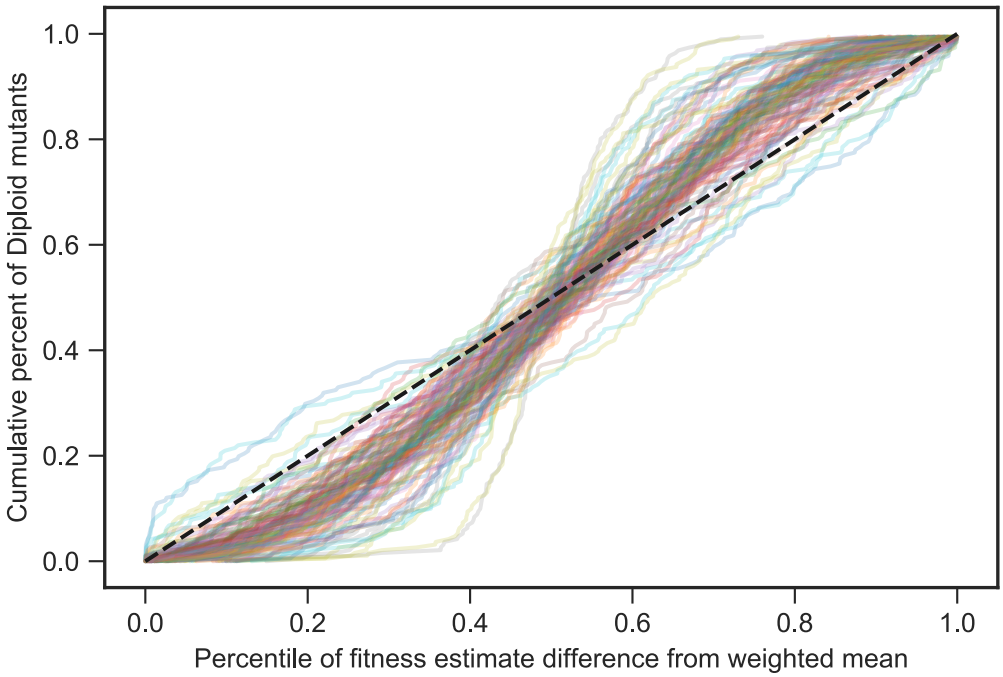

Fig S1. Noise model is a conservative measure of uncertainty. Fitness differences among strains that are genetically identical and have very similar fitness effects tell us about the amount of measurement noise. Our strain collection includes 188 diploids that have similar fitnesses and possess no mutations other than diploidy. For each diploid fitness estimate, we calculated the percentile of deviation from the weighted average of all diploid fitness estimates in a particular environment. This is shown on the horizontal axis. The vertical axis shows the cumulative percent of diploids with deviations listed on the horizontal axis. If the noise model perfectly captures the uncertainty of each measurement, then it should be represented by the black dashed line, as, for instance, 20% of the diploids should have a difference from the mean in the 20th percentile. Each line represents a single experiment (we have 45 environments each with several replicates for a total of 109 experiments, see Methods). For the vast majority of experiments, the diploids are closer to the mean than predicted by our noise model, as indicated by each line's sigmoidal shape. This indicates that the noise model is conservative.

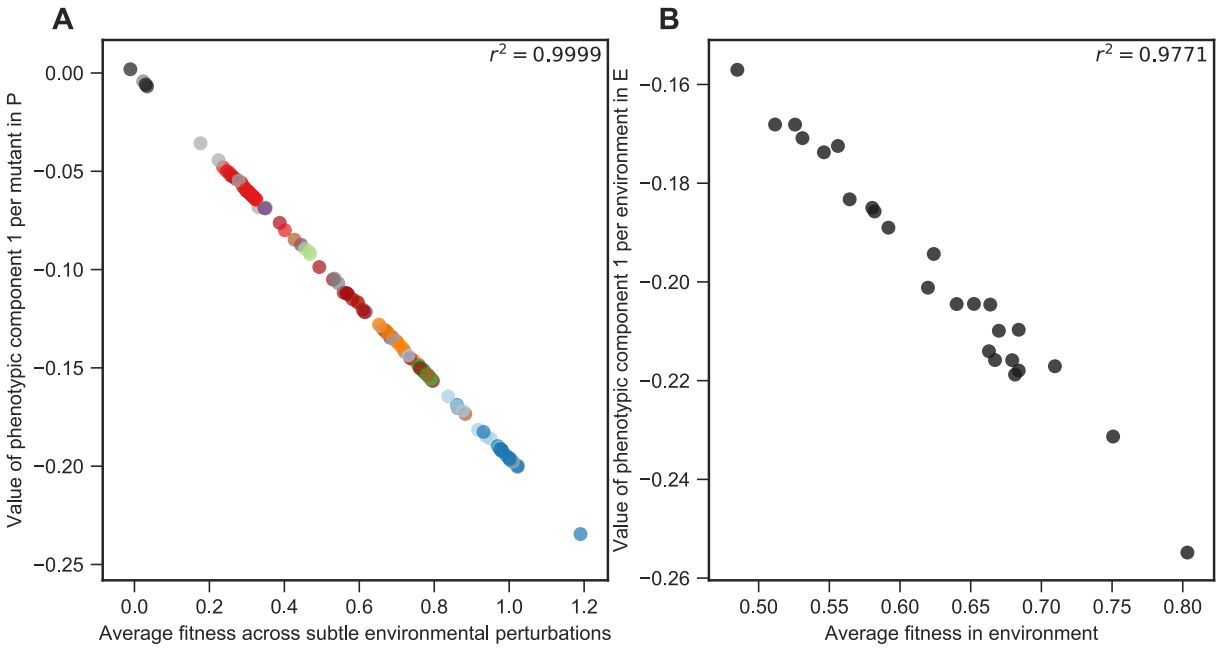

Figure S2. The first component represents the mean fitness of each mutant in the 25 subtle perturbations, as well as the mean impact of each perturbation on fitness. (A) The horizontal axis shows the average fitness of each mutant across all 25 environments that represent subtle perturbations. The vertical axis shows the value of the first phenotypic component for each mutant. Mutants are colored as in Figure 2. **(B)** The horizontal axis shows the average fitness of all 292 mutants in each environment. The vertical axis shows the value of the first phenotypic component in the environment weight space *E.*

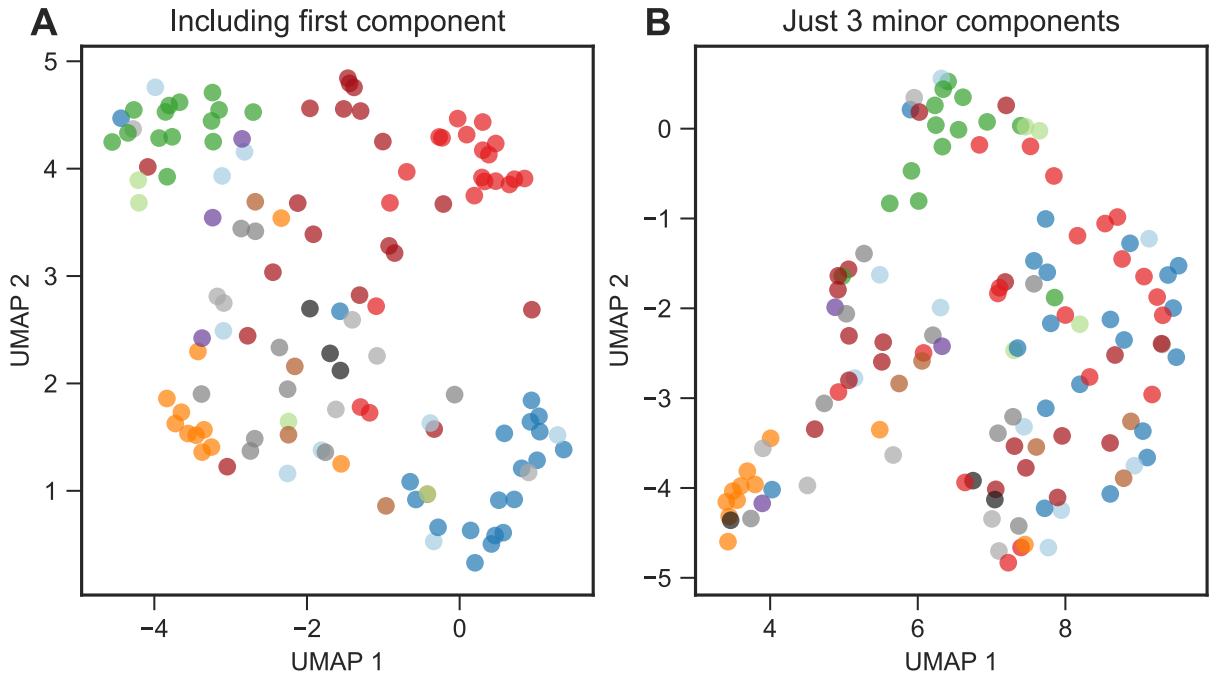

Fig S3. Low-dimensional phenotypic models, and subsets of such models, cluster mutants by gene and mutation type. (A) UMAP clusters mutants visually by gene when using the full 8-component phenotype space. **(B)** UMAP also shows some clustering when using only the three components that explain the least variation in mutant fitness. Though the clustering is clear for *PDE2* and *GPB2*, it less clearly delineates *IRA1-nonsense* and diploid mutants. This suggests these mutants do not have substantial effects on these phenotypic components.

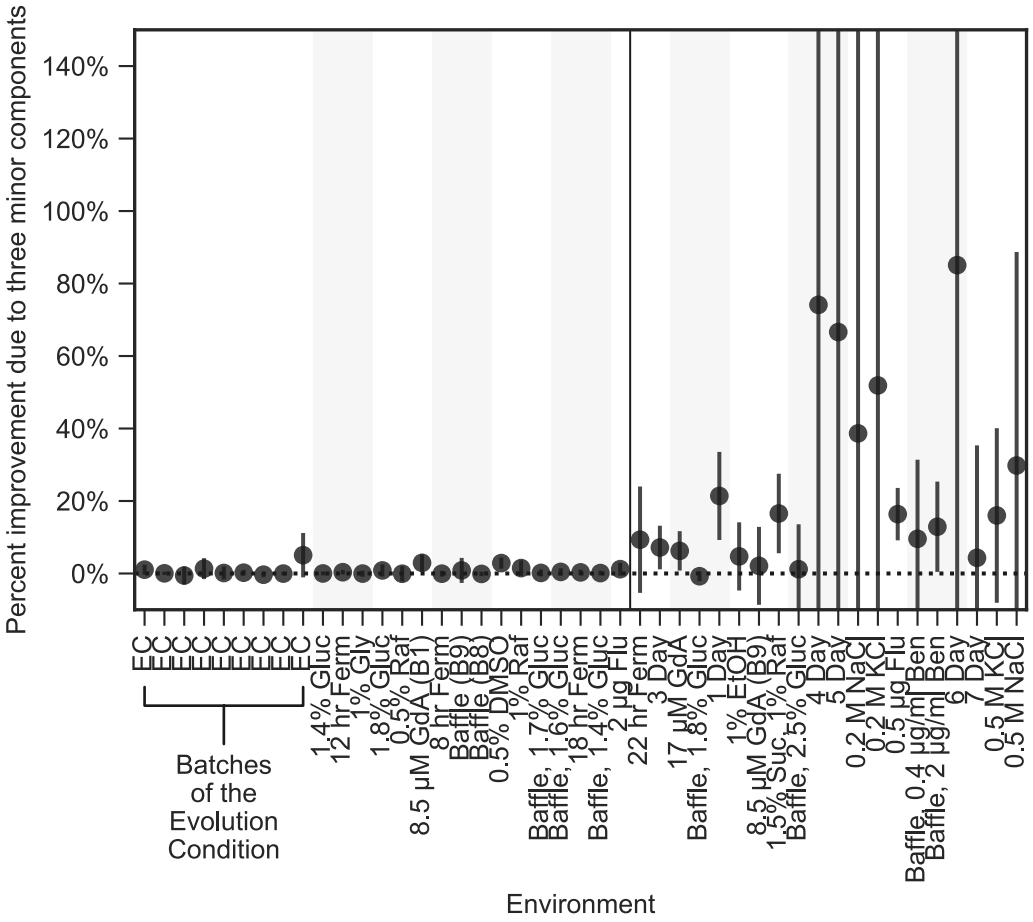

Fig S4. Improved fitness predictions when including the three smallest phenotypic components is not specific to choice of training mutants. This plot is similar to the lower panel of figure 4A, except here, black dots indicate the average improvement across 100 choices of the training and test sets. Error bars indicate two standard deviations from the mean.

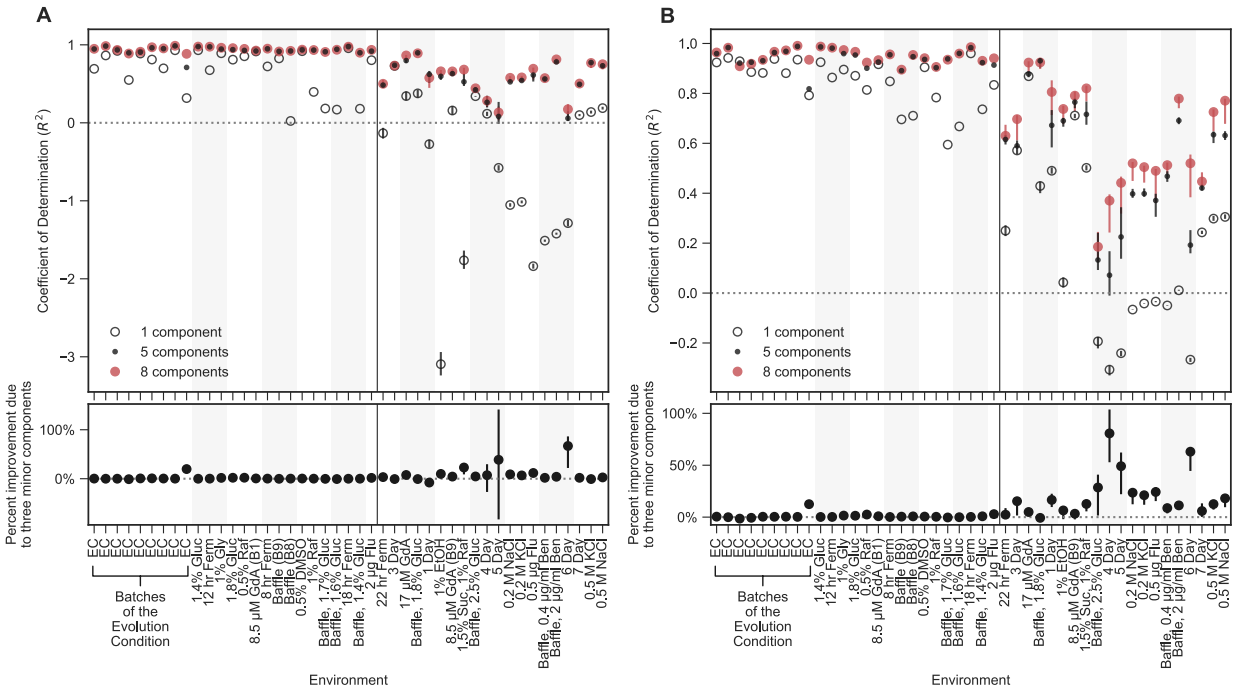

Fig S5. Prediction ability using unweighted coefficient of determination. These plots are similar to figure 4A except here the vertical axis displays prediction power using a standard, rather than a weighted, coefficient of determination measure. Because diploids dominate the number of mutants in the collection, there are large differences between panel A (which shows all mutants) and panel B (which omits diploids).

SUPPLEMENTARY TABLES

Mutation Type	Total Number	Number in Training Set	Number in Test Set
Diploid	188	20	168
High-Fitness Diploid			
Diploid + Chr11Amp	3	1	2
Diploid + Chr12Amp	1	0	1
Diploid + IRA1	1	0	1
Diploid + IRA2	3	1	2
No known add'l mutations	11	5	6
IRA1			
IRA1 nonsense	20	10	10
IRA1 missense	9	4	5
IRA1 other	1	0	1
IRA2	8	4	4
GPB1	4	2	2
GPB2	14	7	7
PDE2	11	5	6
Other Ras/PKA pathway			
CYR1	3	1	2
TFS1	1	0	1
RAS2	1	0	1
TOR/Sch9 pathway			
KOG1	1	0	1
SCH9	1	0	1
TOR1	1	0	1
Other Adaptive	7	0	7
Neutral	3	0	3
TOTAL	**292**	**60**	**232**

Table S1. List of all mutants included in this study.

#	Shortname	Source	Number of barcoded clones	Batch Number	Number of replicates	Description of Manipulation
1	EC	This study	500	10	3	No preculture
2	EC	This study	4800	4	3	No barcodeless ancestor
3	EC	Li et al (2018)	4800	4	3	
4	EC	This study	500	5	3	No preculture
5	EC	Venkataram et al (2016)	4800	1	3	
6	EC	Li et al (2018)	4800	3	3	
7	EC	Venkataram et al (2016)	4800	2	3	
8	EC	This study	500	9	4	
9	EC	This study	500	5	3	
10	1.4% Gluc	This study	500	9	2	1.4% glucose concentration
11	12hr Ferm	Li et al (2018)	4800	3	3	44 hours of growth, 2*10^8 cells transferred, resulting in ~12 hours of fermentation phase
12	1% Gly	This study	500	9	2	Added 1% glycerol
13	1.8% Gluc	This study	500	9	2	1.8% glucose concentration
14	0.5% Raf	This study	500	9	2	Added 0.5% raffinose
15	8.5uM GdA (B1)	This study	4800	1	3	Added 8.5uM Geldanamycin
16	8hr Ferm	Li et al (2018)	4800	3	3	40 hours of growth, 8*10^8 cells transferred, resulting in ~8 hours of fermentation phase
17	Baffle (B9)	This study	500	9	2	Used baffled flask
18	Baffle (B8)	This study	4800	8	2	Used baffled flask
19	0.5% DMSO	This study	4800	1	3	Included 0.5% DMSO
20	1% Raf	This study	500	9	2	Added 1% raffinose
21	Baffle, 1.7% Gluc	This study	4800	8	2	1.7% glucose concentration, used baffled flask
22	Baffle, 1.6% Gluc	This study	4800	8	2	1.6% glucose concentration, used baffled flask
23	18hr Ferm	Li et al (2018)	4800	4	3	50 hours of growth, 2.5*10^7 cells transferred, resulting in ~18 hours of fermentation phase
24	Baffle, 1.4% Gluc	This study	4800	8	2	1.4% glucose concentration, used baffled flask
25	2ug Flu	This study	500	9	2	Added 2ug Fluconazole
26	22hr Ferm	Li et al (2018)	4800	4	3	54 hours of growth, 6.25*10^6 cells transferred, resulting in ~22 hours of fermentation phase
27	3 Day	Li et al (2018)	4800	2	3	3 days of growth
28	17uM GdA	This study	500	9	2	Added 17uM Geldanamycin
29	Baffle, 1.8% Gluc	This study	4800	8	2	1.8% glucose concentration, used baffled flask
30	1 Day	Li et al (2018)	4800	2	3	24 hours of growth
31	1% EtOH	This study	500	9	2	Added 1% ethanol
32	8.5uM GdA (B9)	This study	500	9	2	Added 8.5uM Geldanamycin
33	1.5% Suc, 1% Raf	This study	500	9	2	No glucose, 1.5% Sucrose, 1% Raffinose
34	Baffle, 2.5% Gluc	This study	4800	8	2	2.5% glucose concentration, used baffled flask
35	4 Day	Li et al (2018)	4800	2	3	4 days of growth
36	5 Day	Li et al (2018)	4800	6	3	5 days of growth
37	0.2M NaCl	This study	500	9	2	Added 0.2M NaCl
38	0.2M KCl	This study	500	9	1	Added 0.2M KCl
39	0.5ug Flu	This study	500	9	2	Added 0.5ug Fluconazole
40	Baffle, 0.4ug/ml Ben	This study	4800	7	2	Added 0.4ug/mL Benomyl, used baffled flask
41	Baffle, 2ug/ml Ben	This study	4800	7	2	Added 2ug/mL Benomyl, used baffled flask
42	6 Day	Li et al (2018)	4800	6	3	6 days of growth
43	7 Day	Li et al (2018)	4800	6	3	7 days of growth
44	0.5M KCl	This study	500	9	1	Added 0.5M KCl
45	0.5M NaCl	This study	500	9	1*	Added 0.5M NaCl

Table S2. List of all conditions used in this study, ordered by deviation from the EC batch as in the main text figures.